

O'REILLY®

oscon

Scalable graph analysis with Apache Giraph and Spark GraphX

oscon.com
#oscon

Roman Shaposhnik rvs@apache.org @rhatr
Director of Open Source, Pivotal Inc.

O'REILLY®

oscon

Introduction into scalable graph analysis with Apache Giraph and Spark GraphX

oscon.com
#oscon

Roman Shaposhnik rvs@apache.org @rhatr
Director of Open Source, Pivotal Inc.

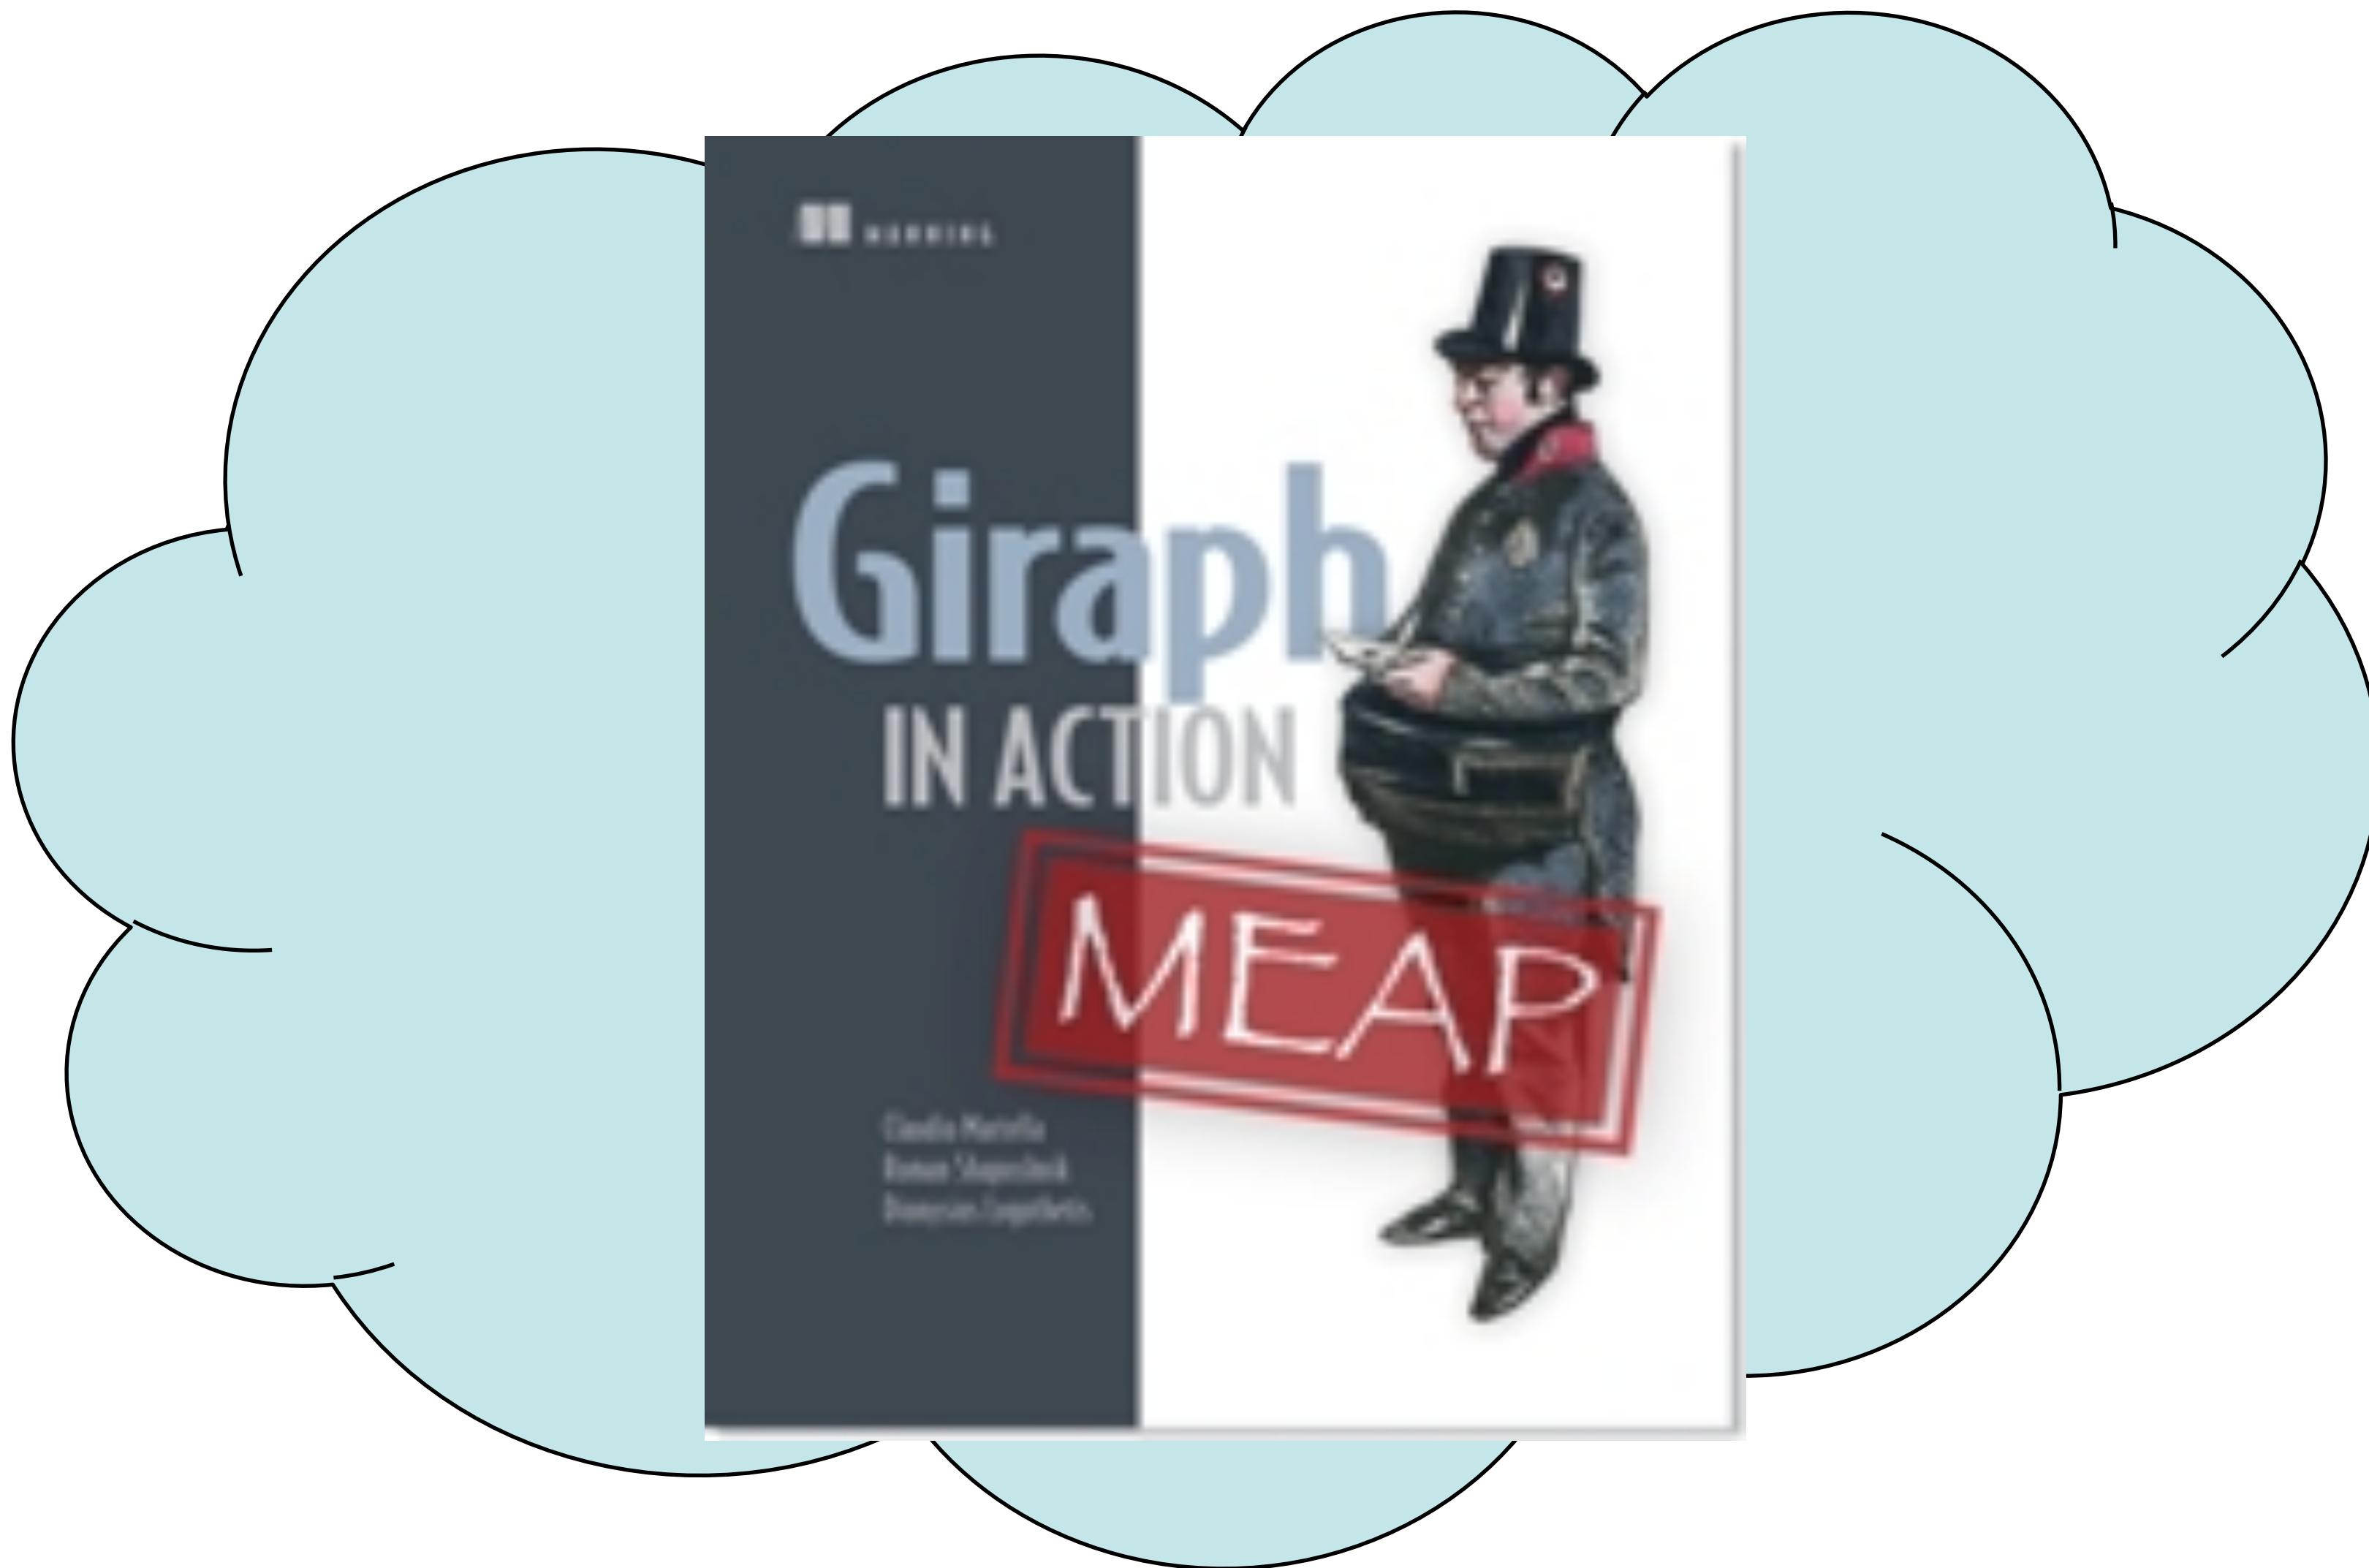
Shameless plug #1



#oscon

O'REILLY®
OSCON

Shameless plug #1



#oscon

O'REILLY®
OSCON

Agenda:

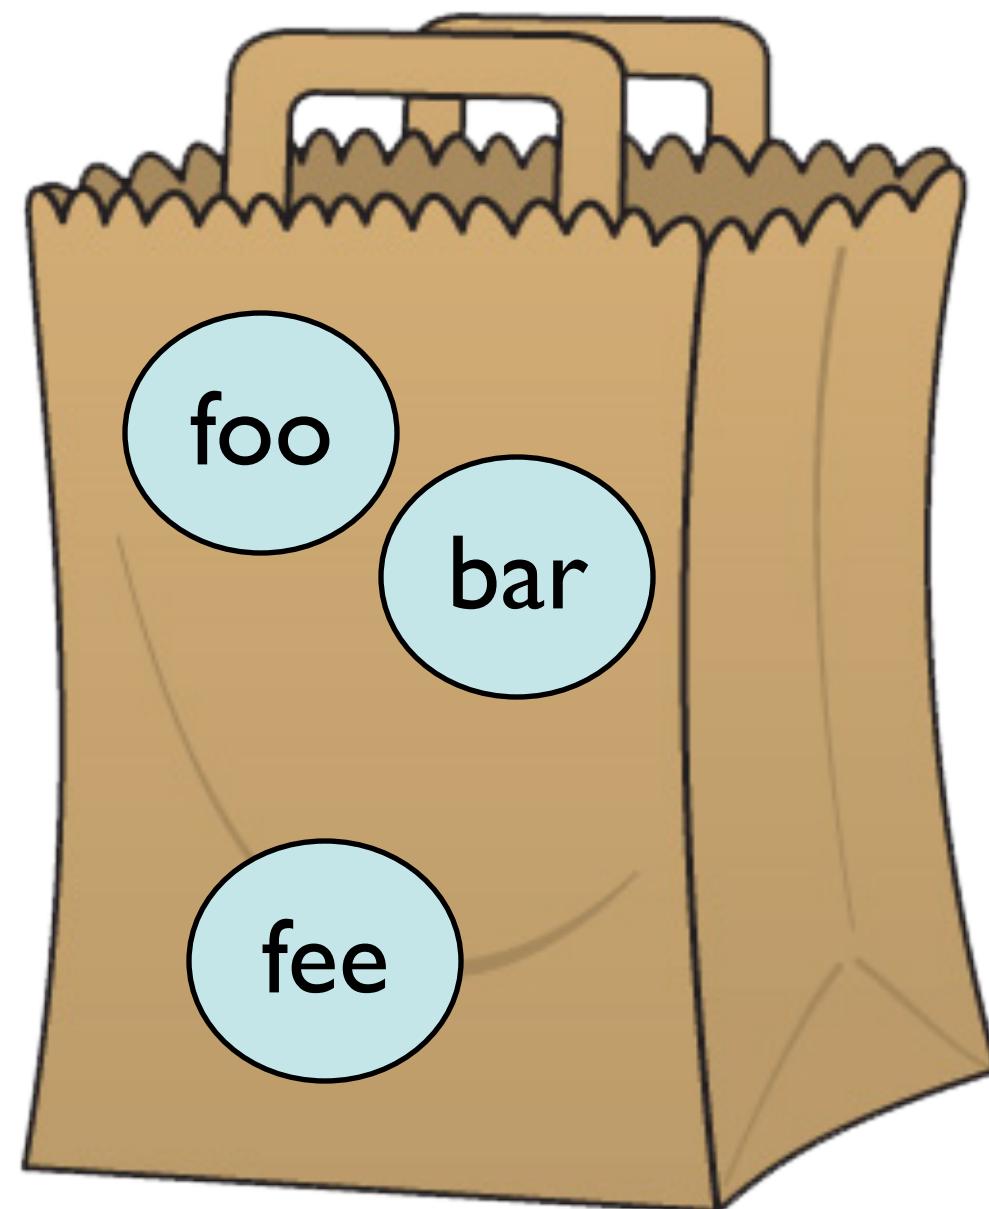


#oscon

O'REILLY®
OSCON

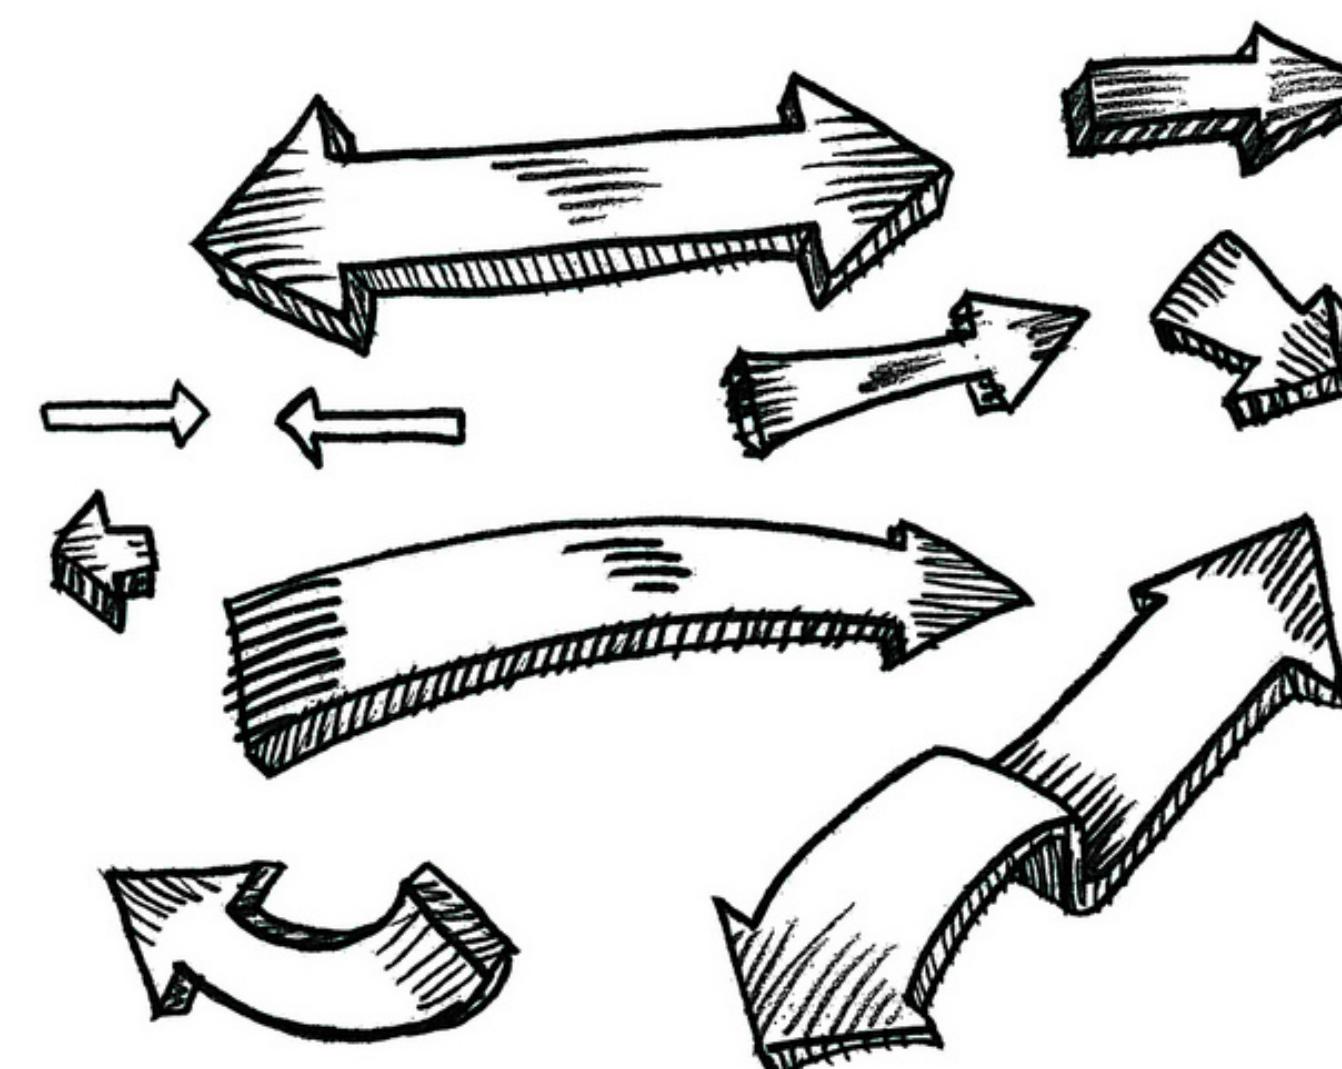
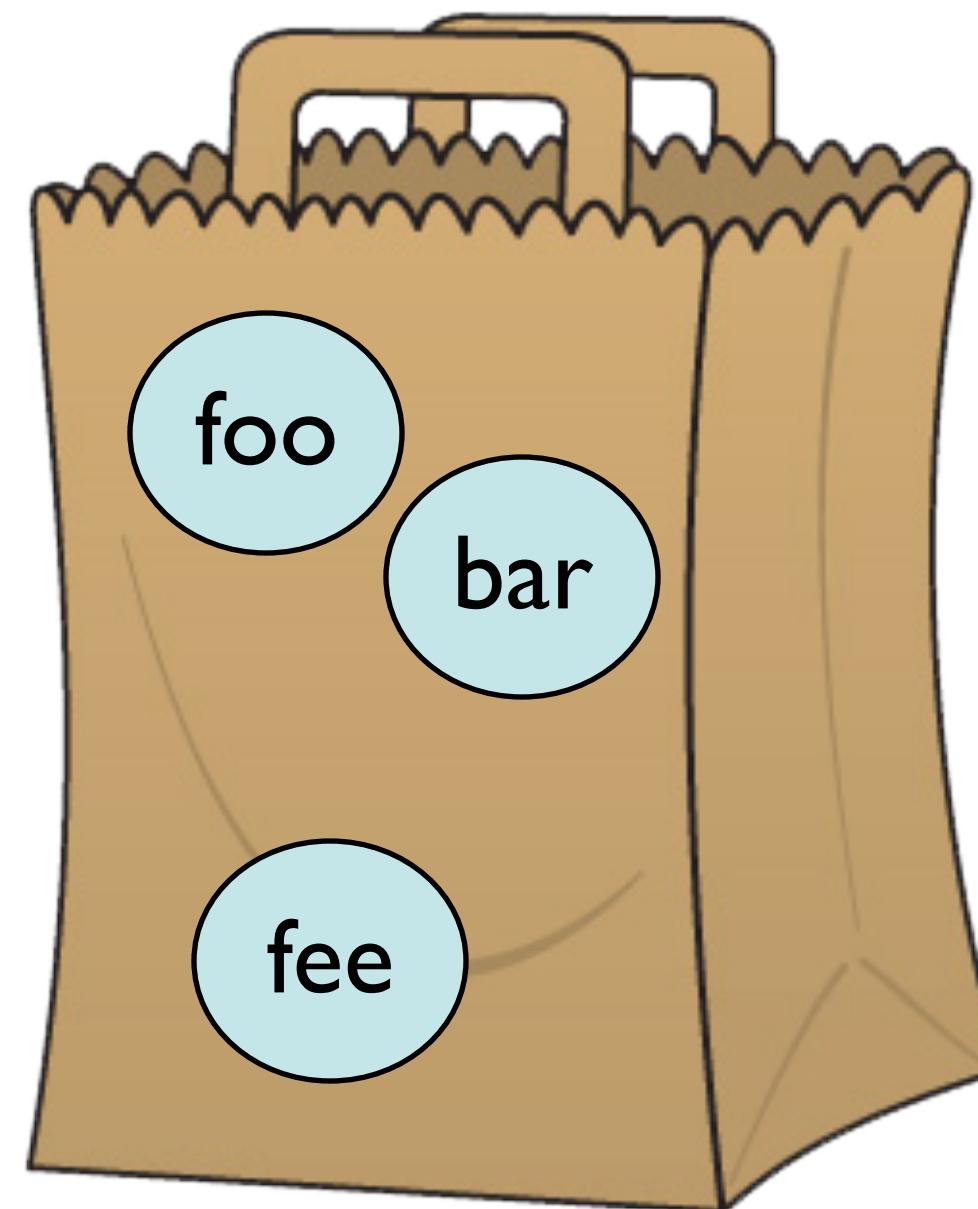
Lets define some terms

- Graph is a $G = (V, E)$, where $E \subseteq V \times V$
- Directed multigraphs with properties attached to each vertex and edge



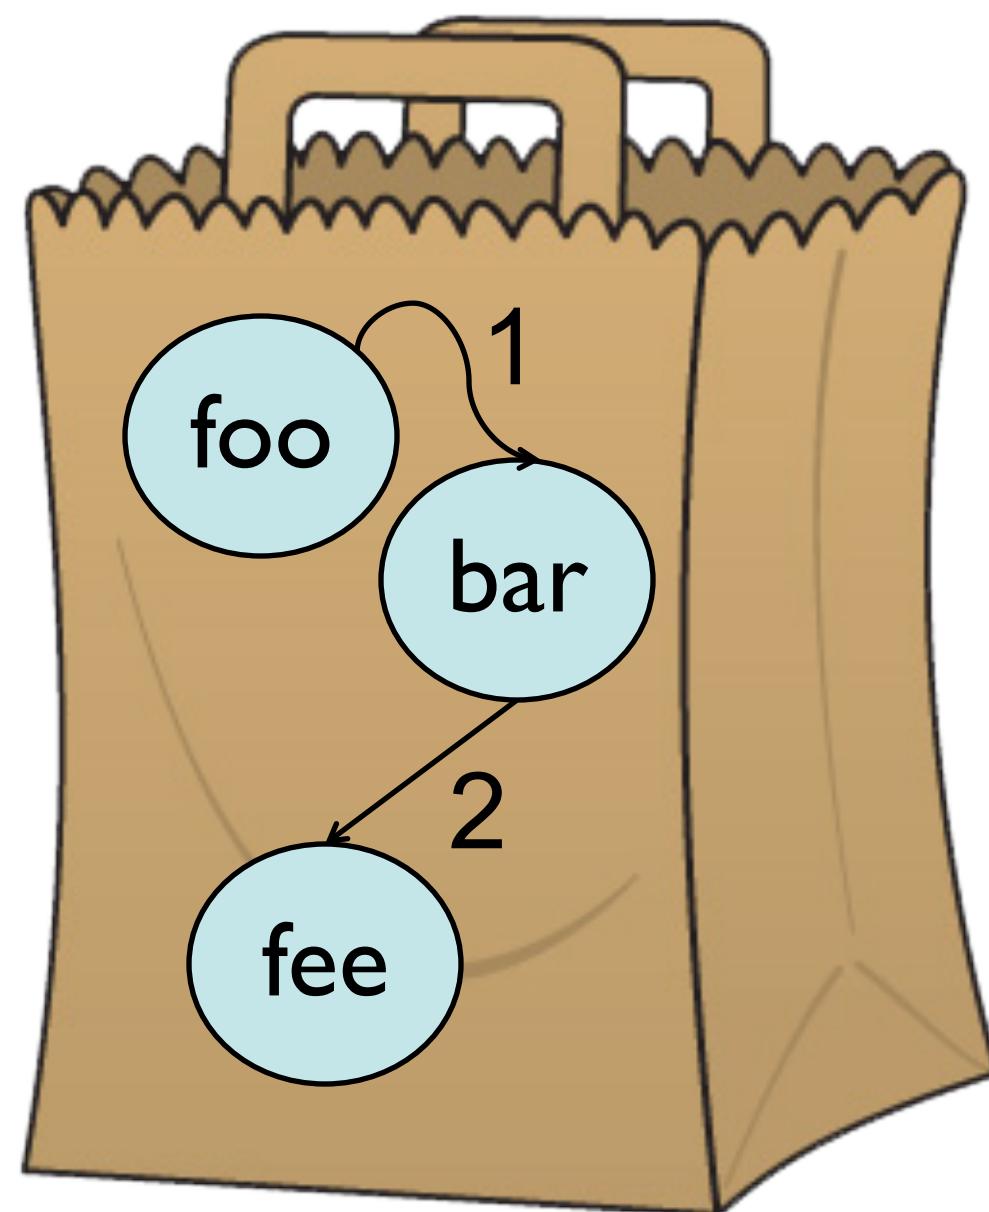
Lets define some terms

- Graph is a $G = (V, E)$, where $E \subseteq V \times V$
- Directed multigraphs with properties attached to each vertex and edge



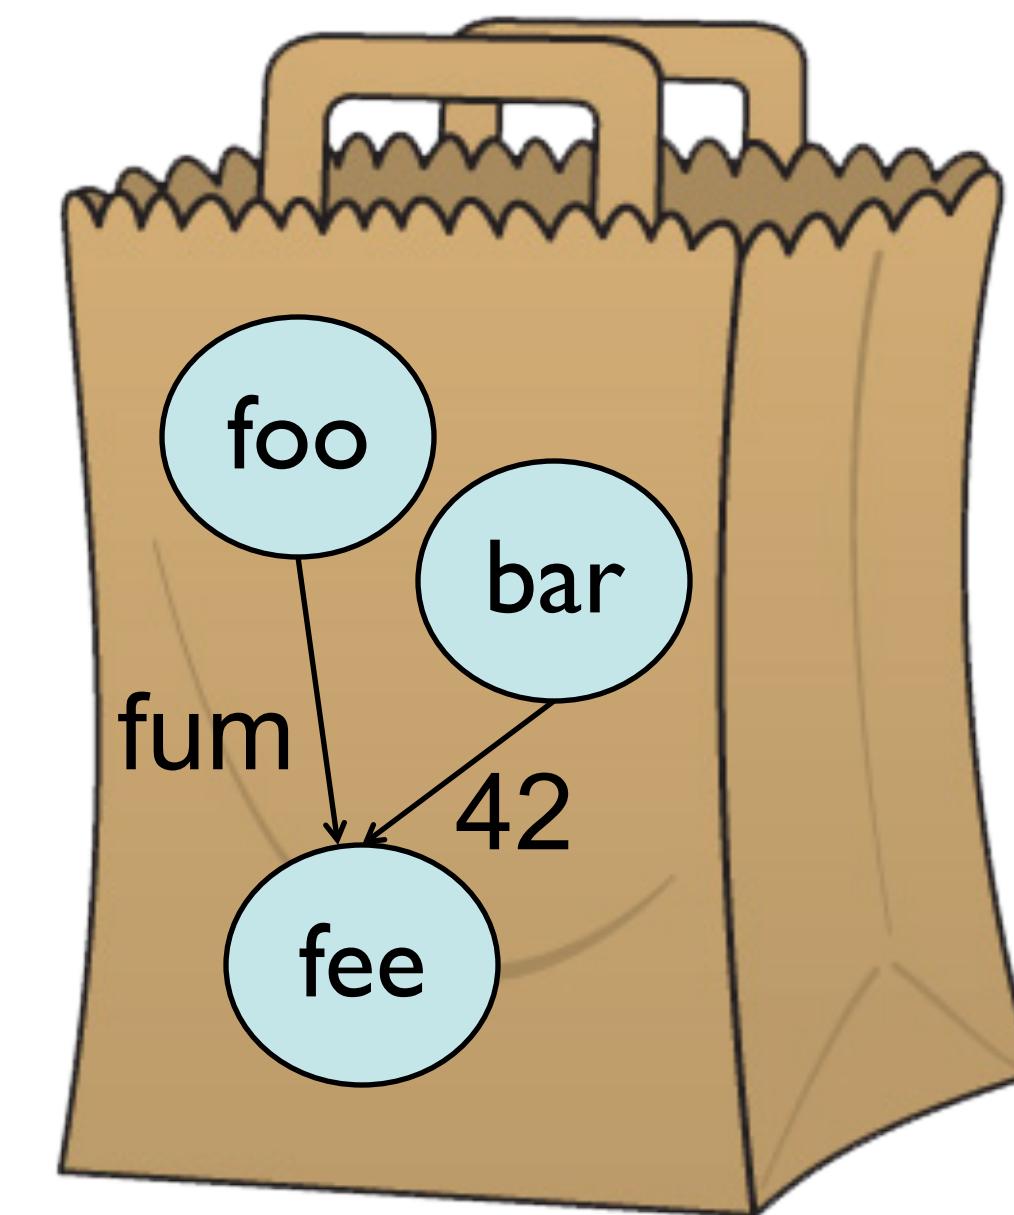
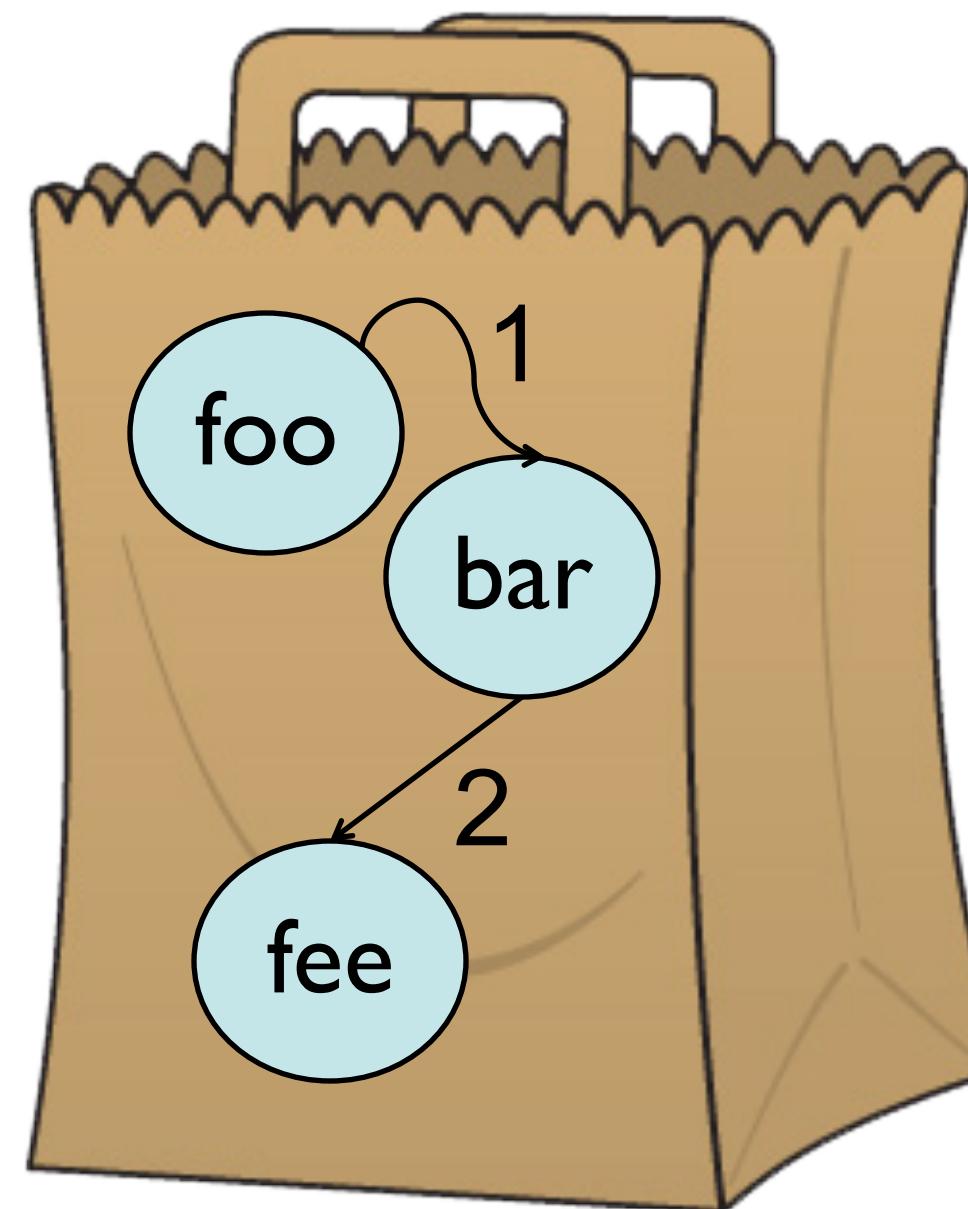
Lets define some terms

- Graph is a $G = (V, E)$, where $E \subseteq V \times V$
- Directed multigraphs with properties attached to each vertex and edge



Lets define some terms

- Graph is a $G = (V, E)$, where $E \subseteq V \times V$
- Directed multigraphs with properties attached to each vertex and edge



What kind of graphs are we talking about?

- Page ranking on Facebook social graph (mid 2013)
 - 10^9 (billions) vertices
 - 10^{12} (trillion) edges
 - 10^{15} (petabyte) cold storage data scale
 - 200 servers
 - ...all in under 4 minutes!

O'REILLY®

oscon

“On day one Doug created
HDFS and MapReduce”

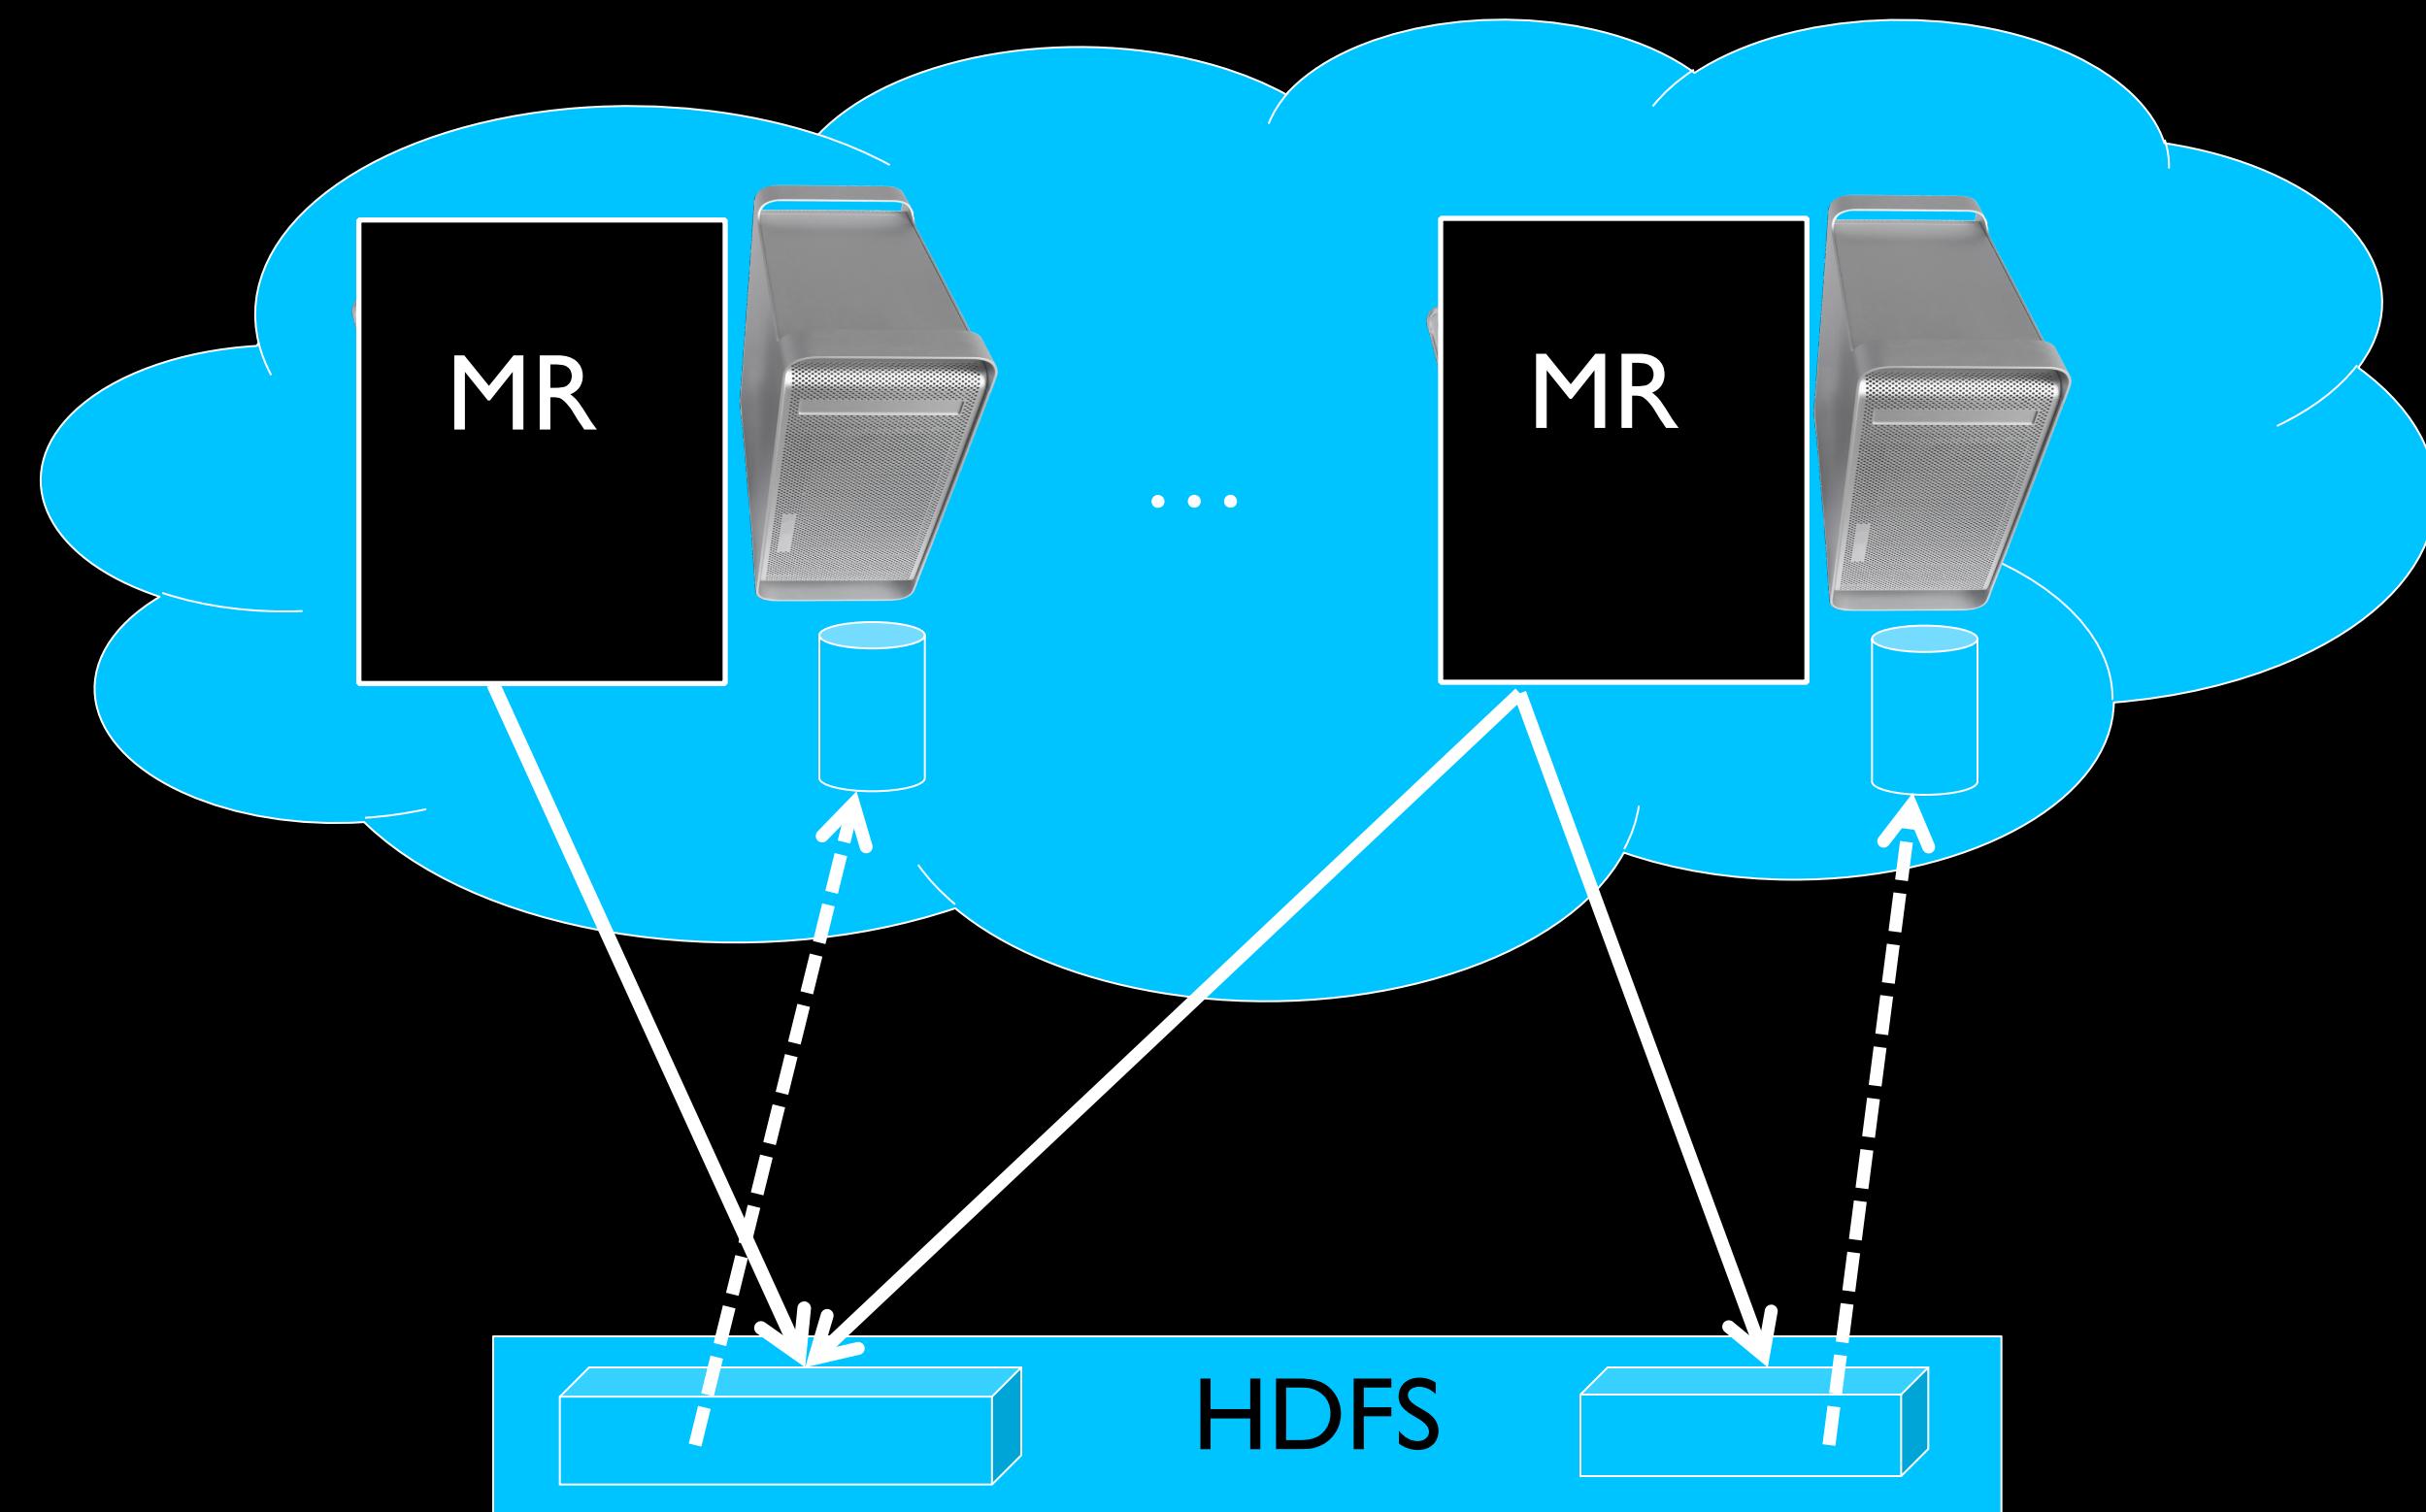
oscon.com

#oscon

Google papers that started it all

- GFS (file system)
 - distributed
 - replicated
 - non-POSIX
- MapReduce (computational framework)
 - distributed
 - batch-oriented (long jobs; final results)
 - data-gravity aware
 - designed for “embarrassingly parallel” algorithms

HDFS pools and abstracts direct-attached storage



A Unix analogy

- It is as though instead of:

```
$ grep foo bar.txt | tr “,” “ “ | sort -u
```

- We are doing:

```
$ grep foo < bar.txt > /tmp/1.txt  
$ tr “,” “ “ < /tmp/1.txt > /tmp/2.txt  
$ sort -u < /tmp/2.txt
```

O'REILLY®

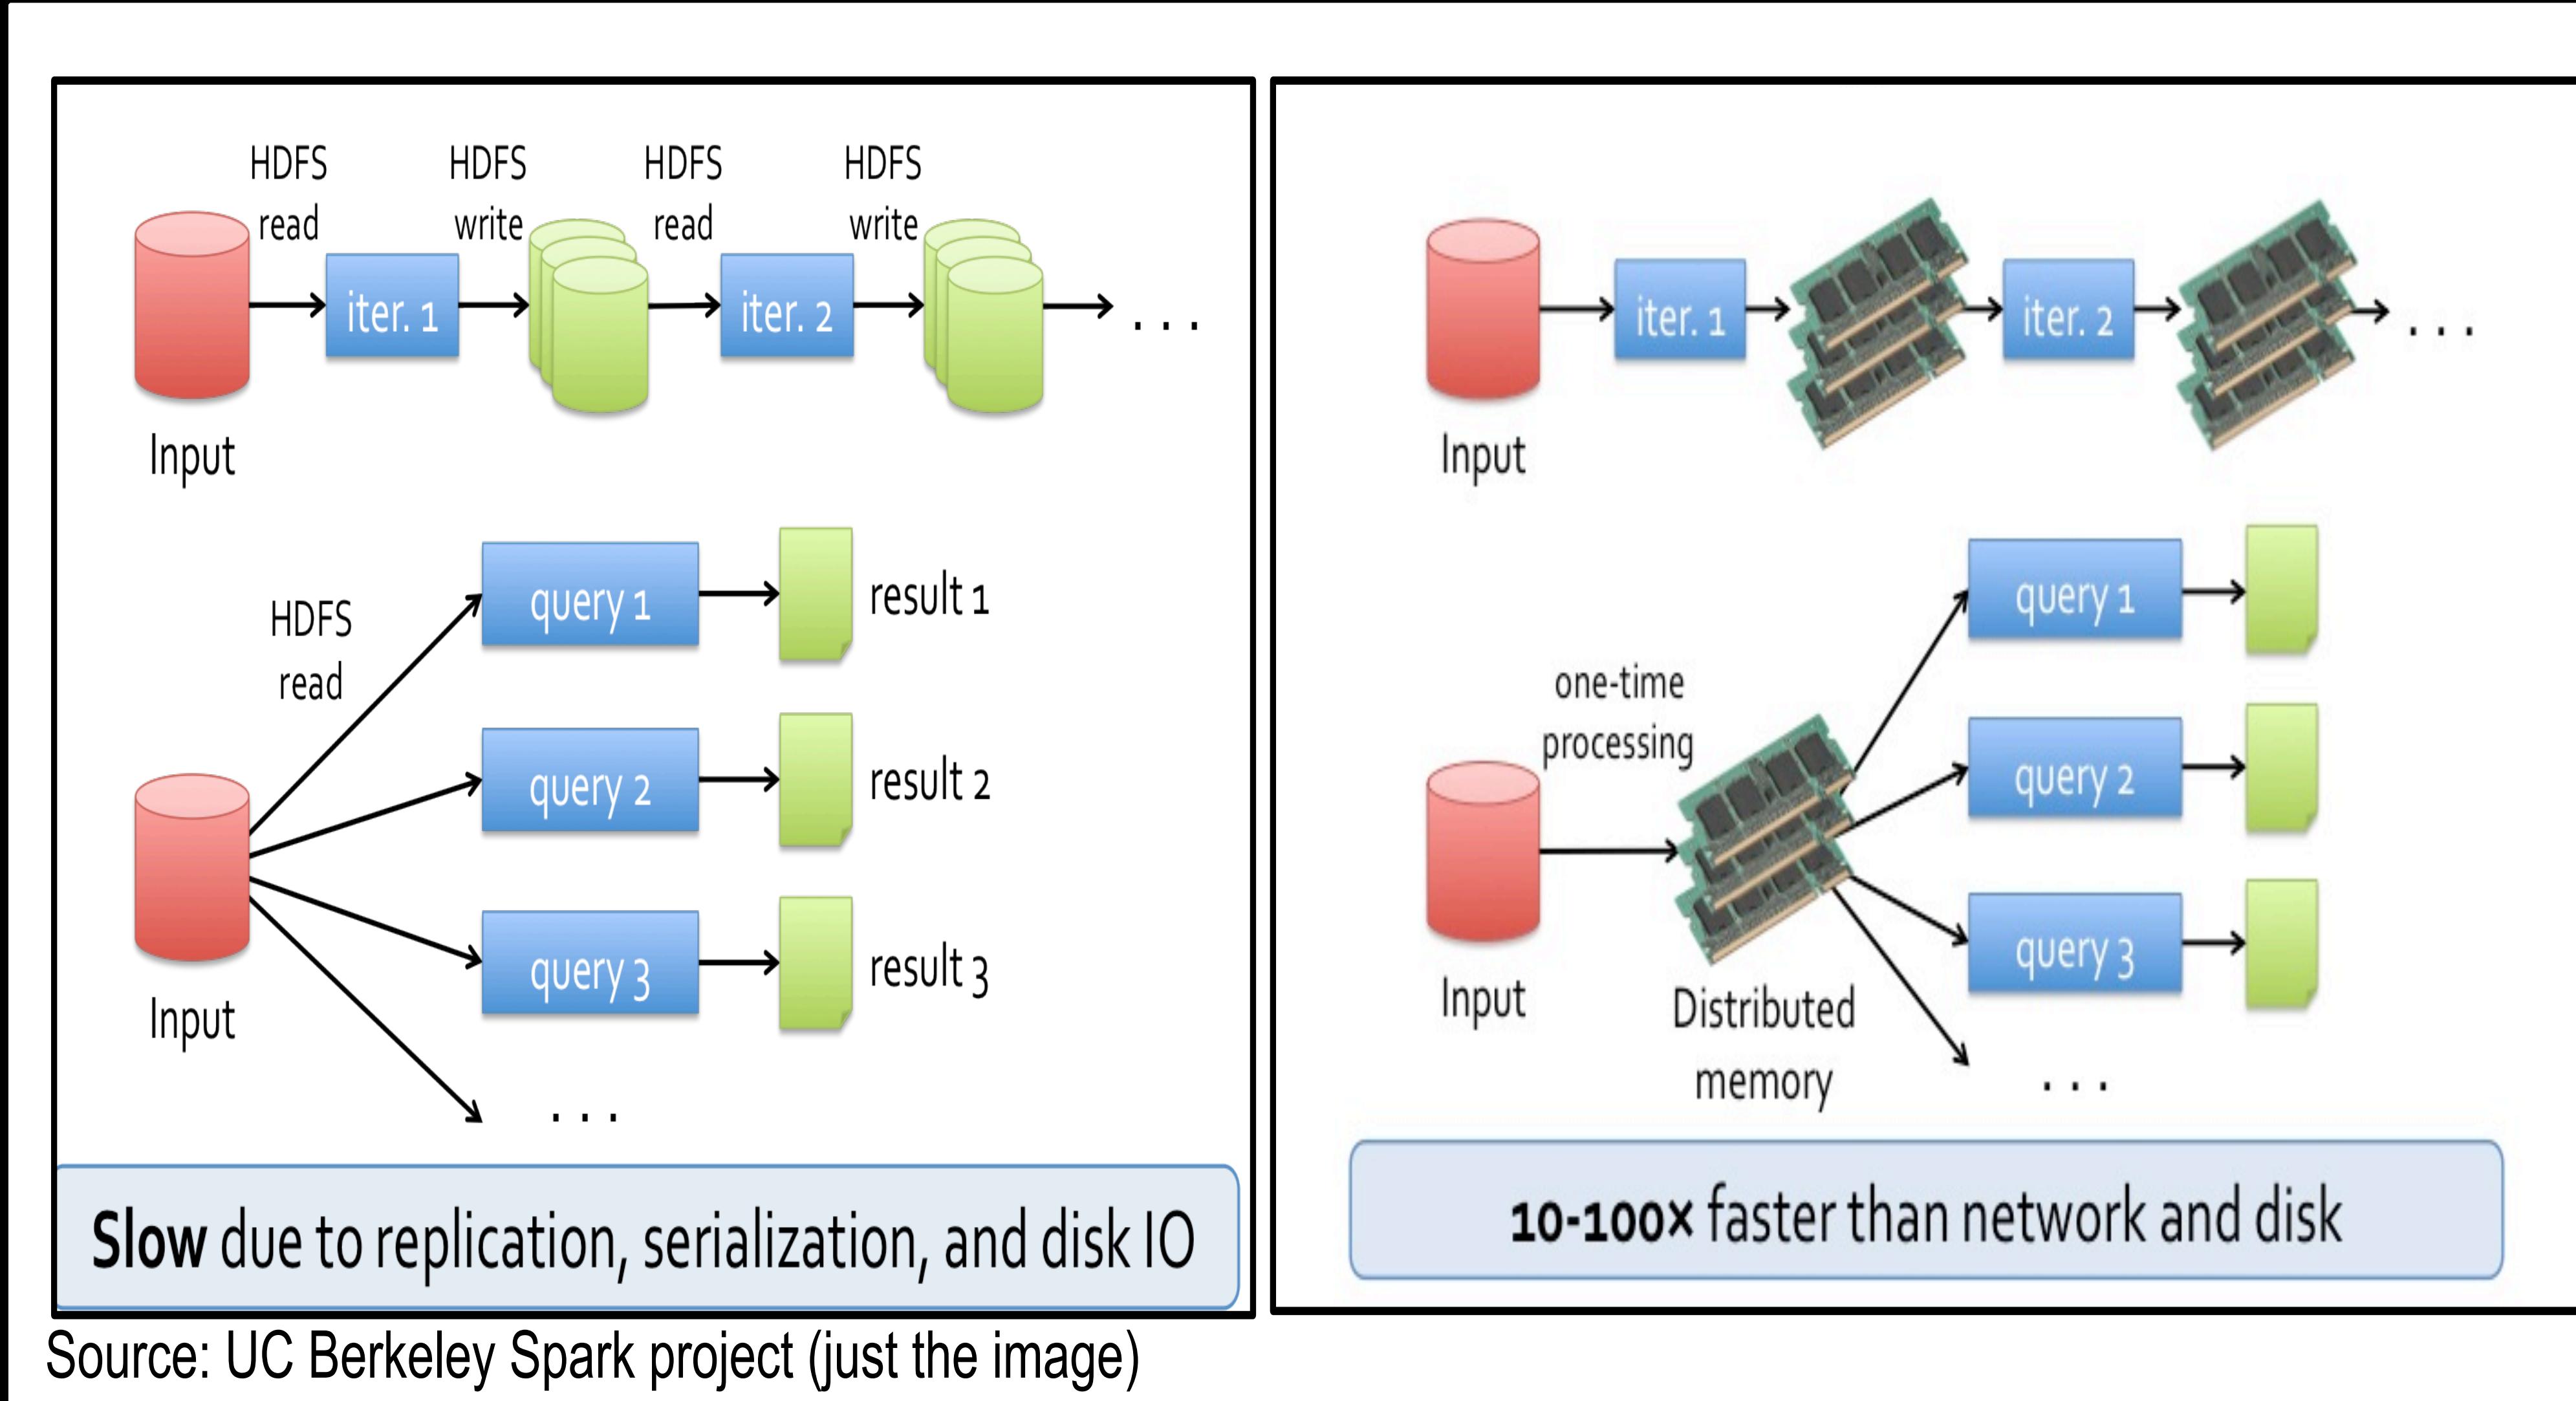
oscon

Enter Apache Spark

oscon.com

#oscon

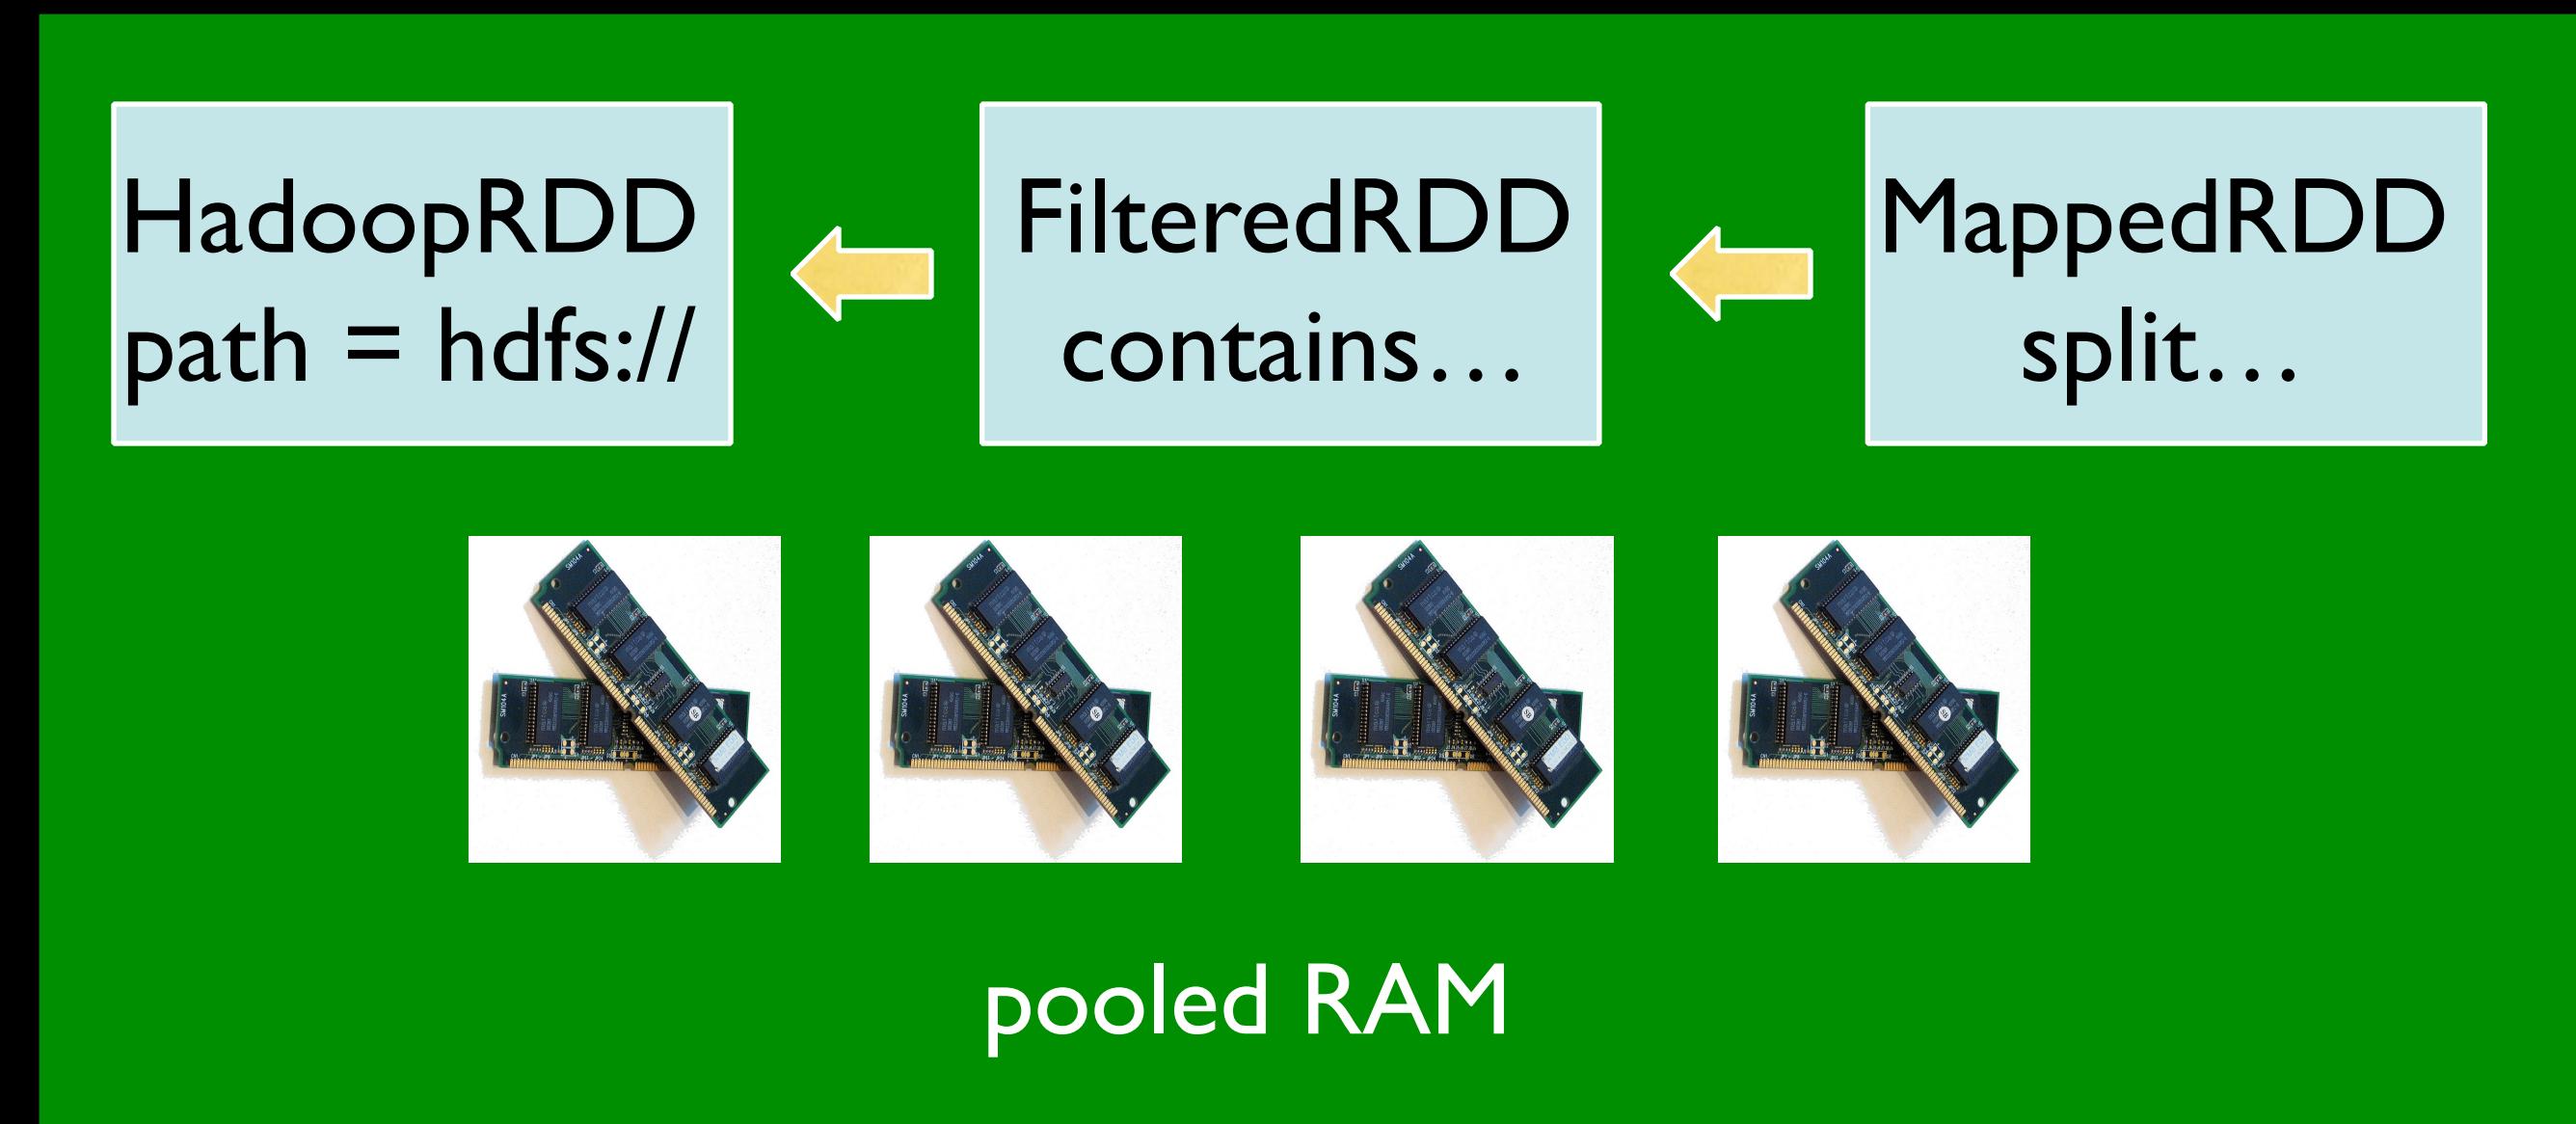
RAM is the new disk, Disk is the new tape



Source: UC Berkeley Spark project (just the image)

RDDs instead of HDFS files, RAM instead of Disk

```
warnings = textFile(...).filter(_.contains("warning"))
    .map(_.split(' ')(1))
```



RDDs: resilient, distributed, datasets

- Distributed on a cluster in RAM
- Immutable (mostly)
- Can be evicted, snapshotted, etc.
- Manipulated via parallel operators (map, etc.)
- Automatically rebuilt on failure
- A parallel ecosystem
- A solution to iterative and multi-stage apps

O'REILLY®

oscon

What's so special about Graphs and
big data?

oscon.com

#oscon

Graph relationships

- Entities in your data: tuples
 - customer data
 - product data
 - interaction data
- Connection between entities: graphs
 - social network or my customers
 - clustering of customers vs. products

A word about Graph databases

- Plenty available
 - Neo4J, Titan, etc.
- Benefits
 - Query language
 - Tightly integrate systems with few moving parts
 - High performance on known data sets
- Shortcomings
 - Not easy to scale horizontally
 - Don't integrate with HDFS
 - Combine storage and computational layers
 - A sea of APIs

What's the key API?

- Directed multi-graph with labels attached to vertices and edges
- Defining vertices and edges dynamically
- Selecting sub-graphs
- Mutating the topology of the graph
- Partitioning the graph
- Computing model that is
 - iterative
 - scalable (shared nothing)
 - resilient
 - easy to manage at scale

O'REILLY®

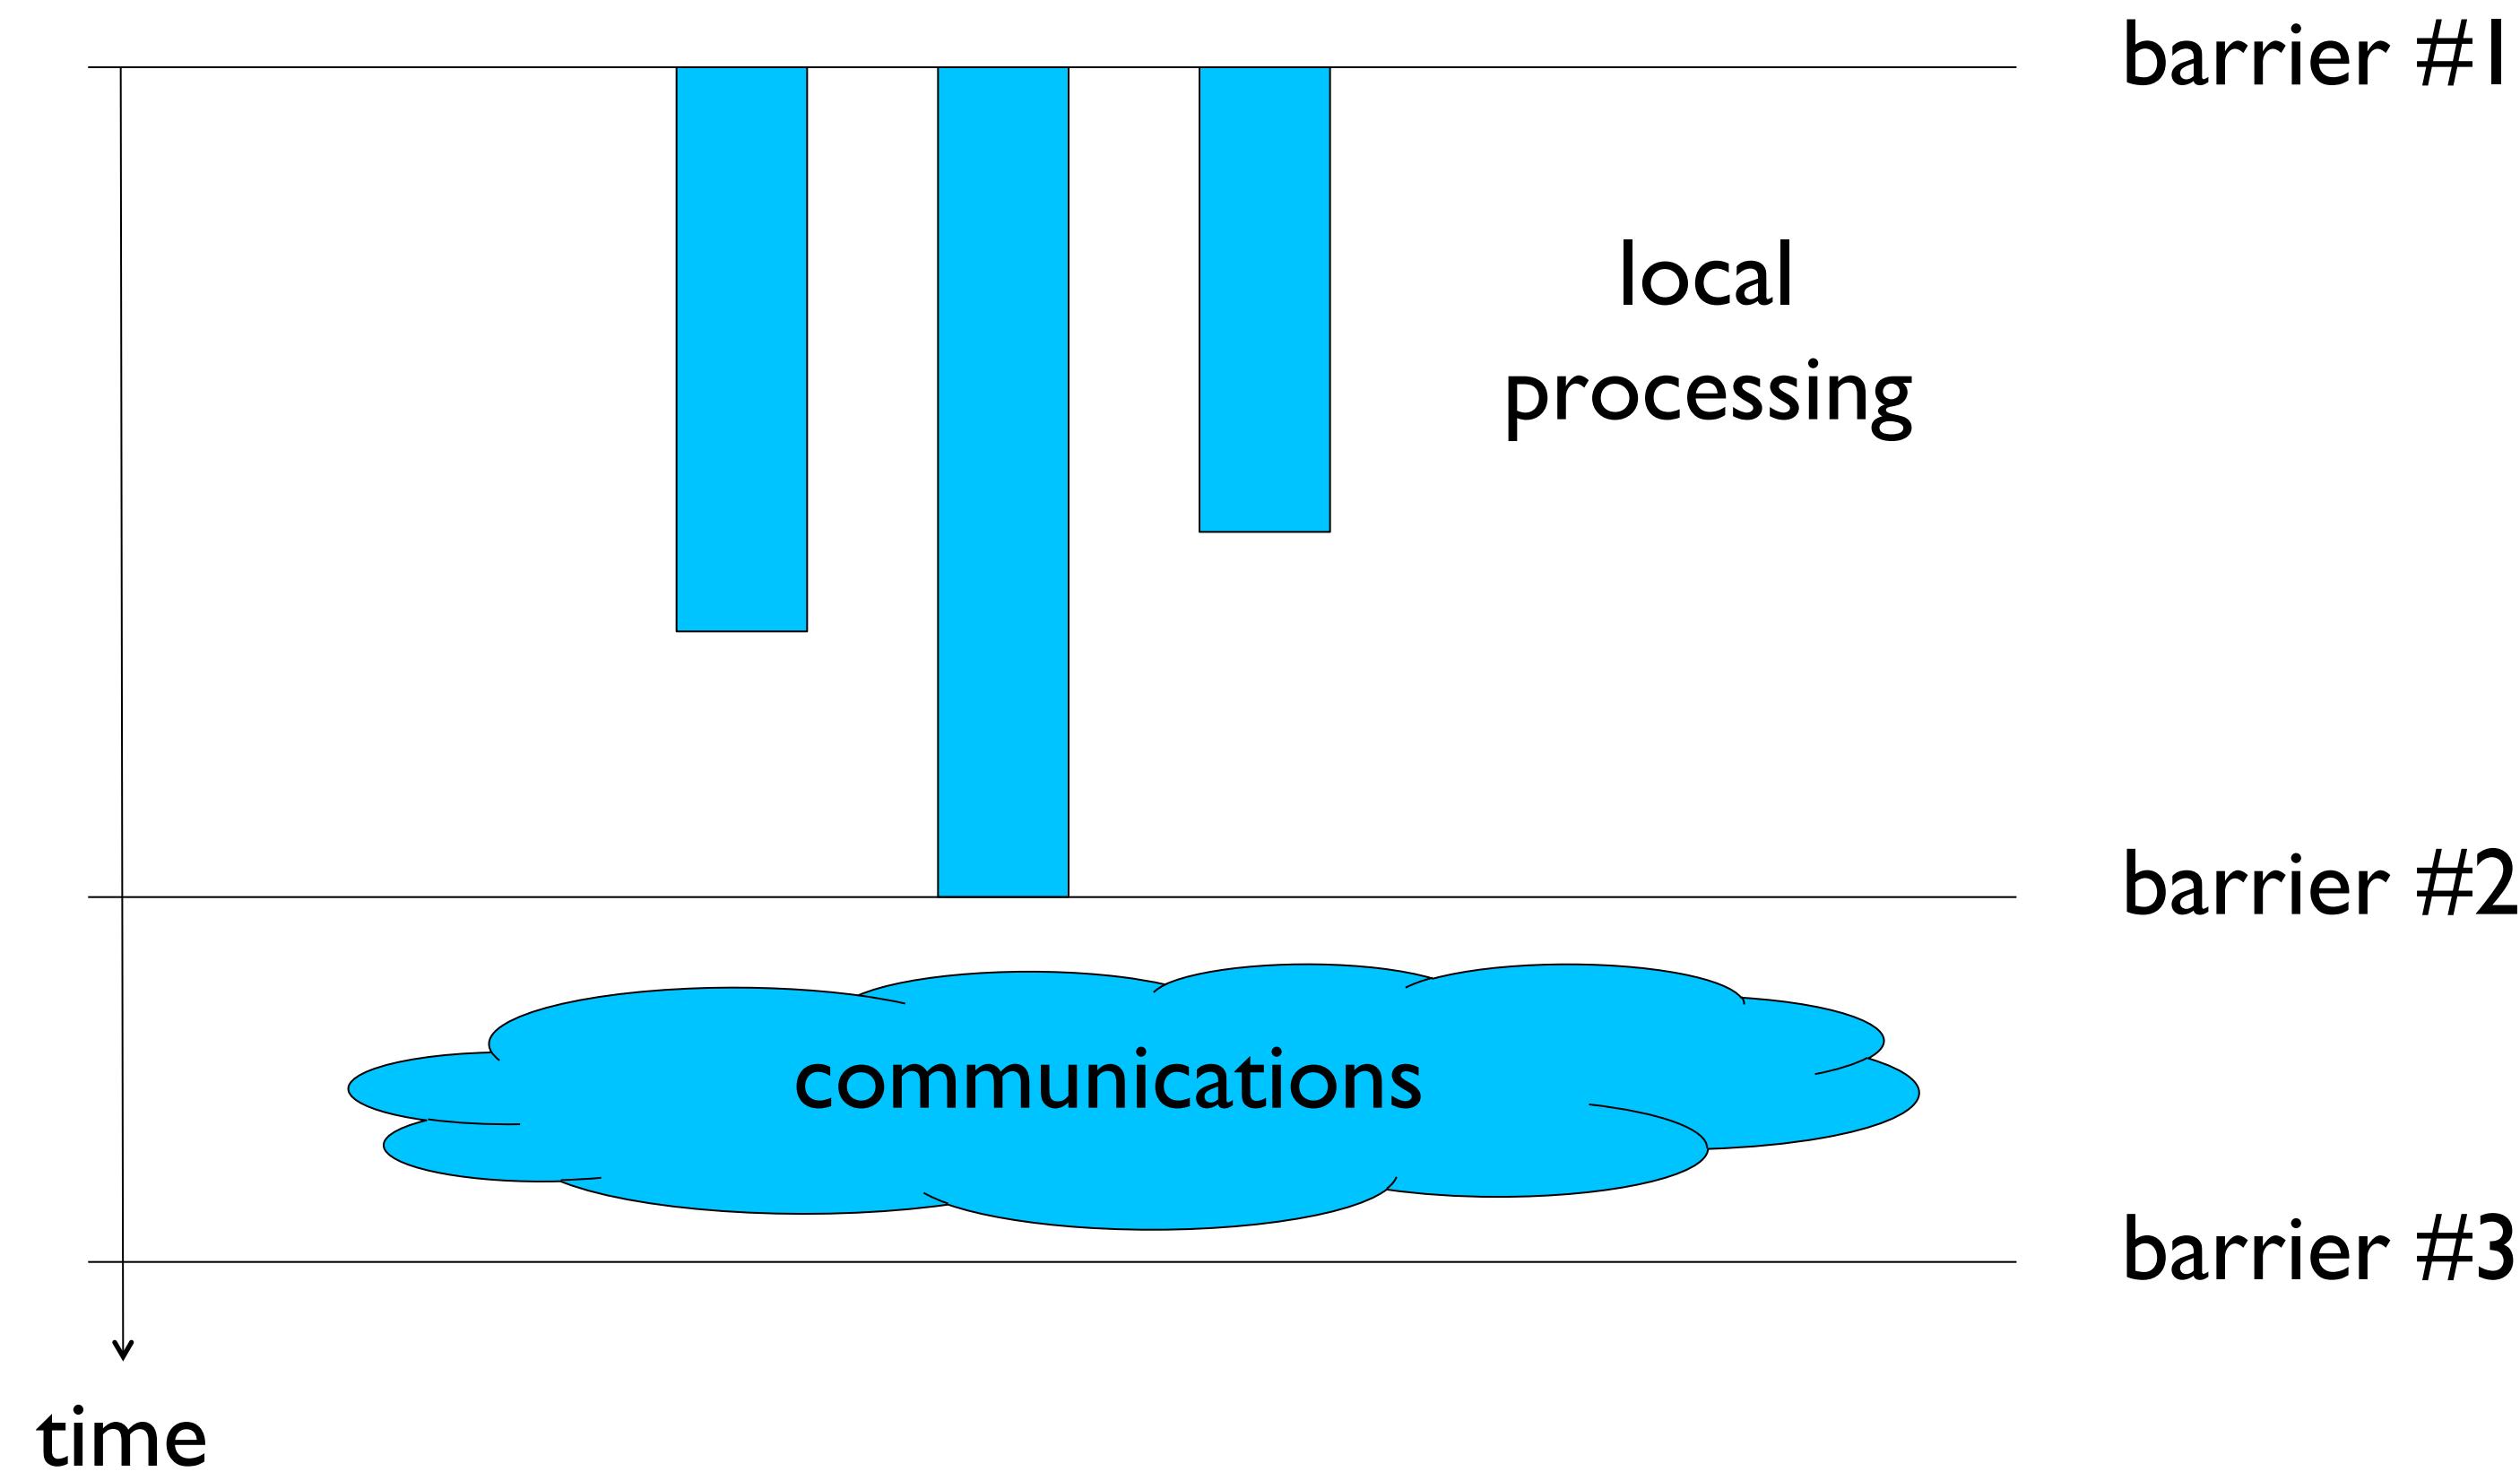
oscon

Bulk Synchronous Parallel BSP compute model

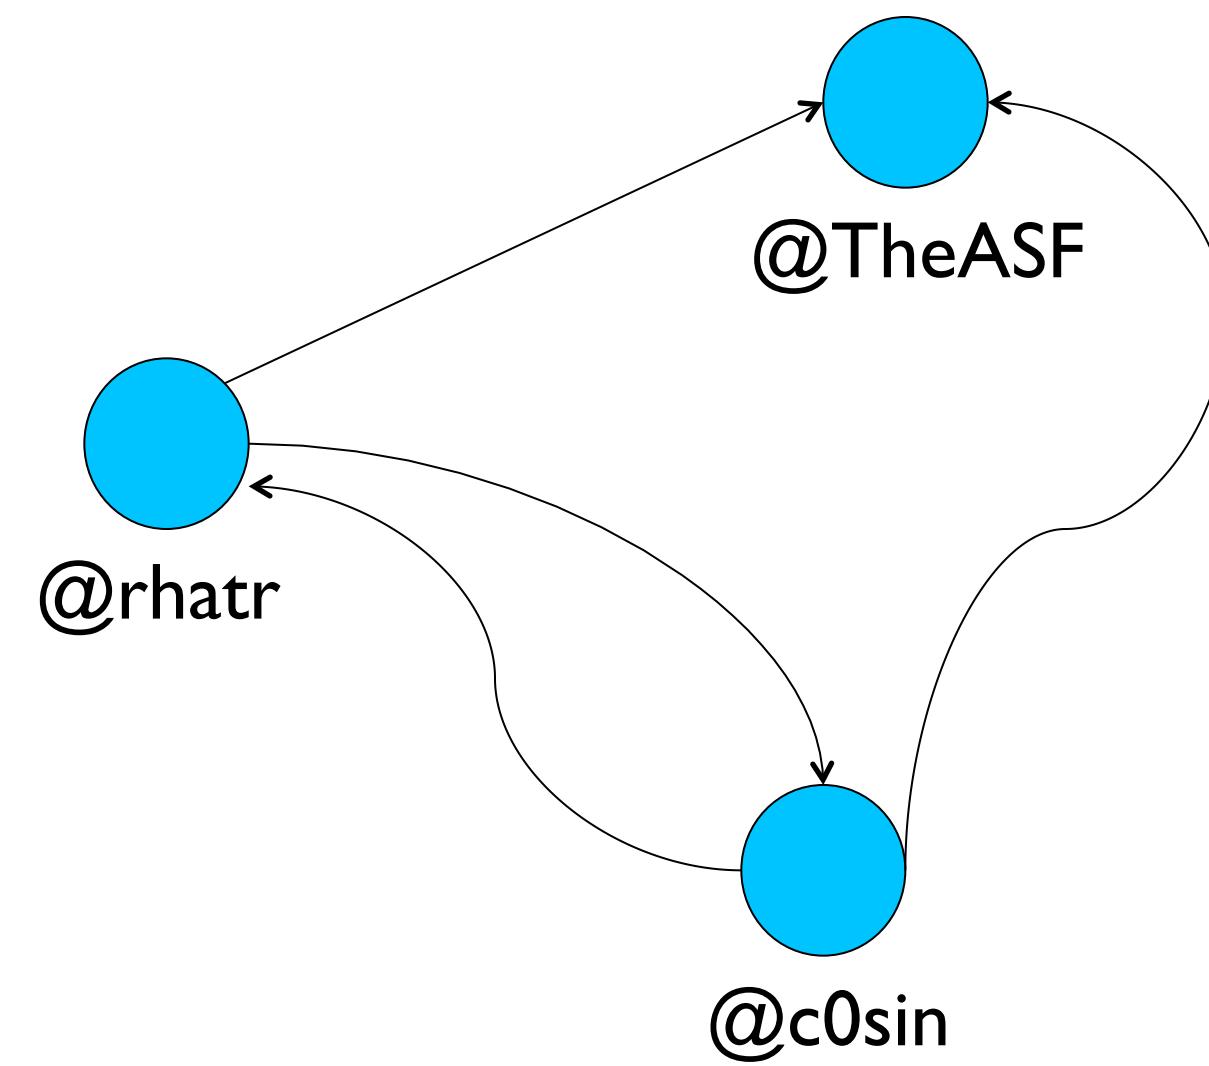
oscon.com

#oscon

BSP in a nutshell



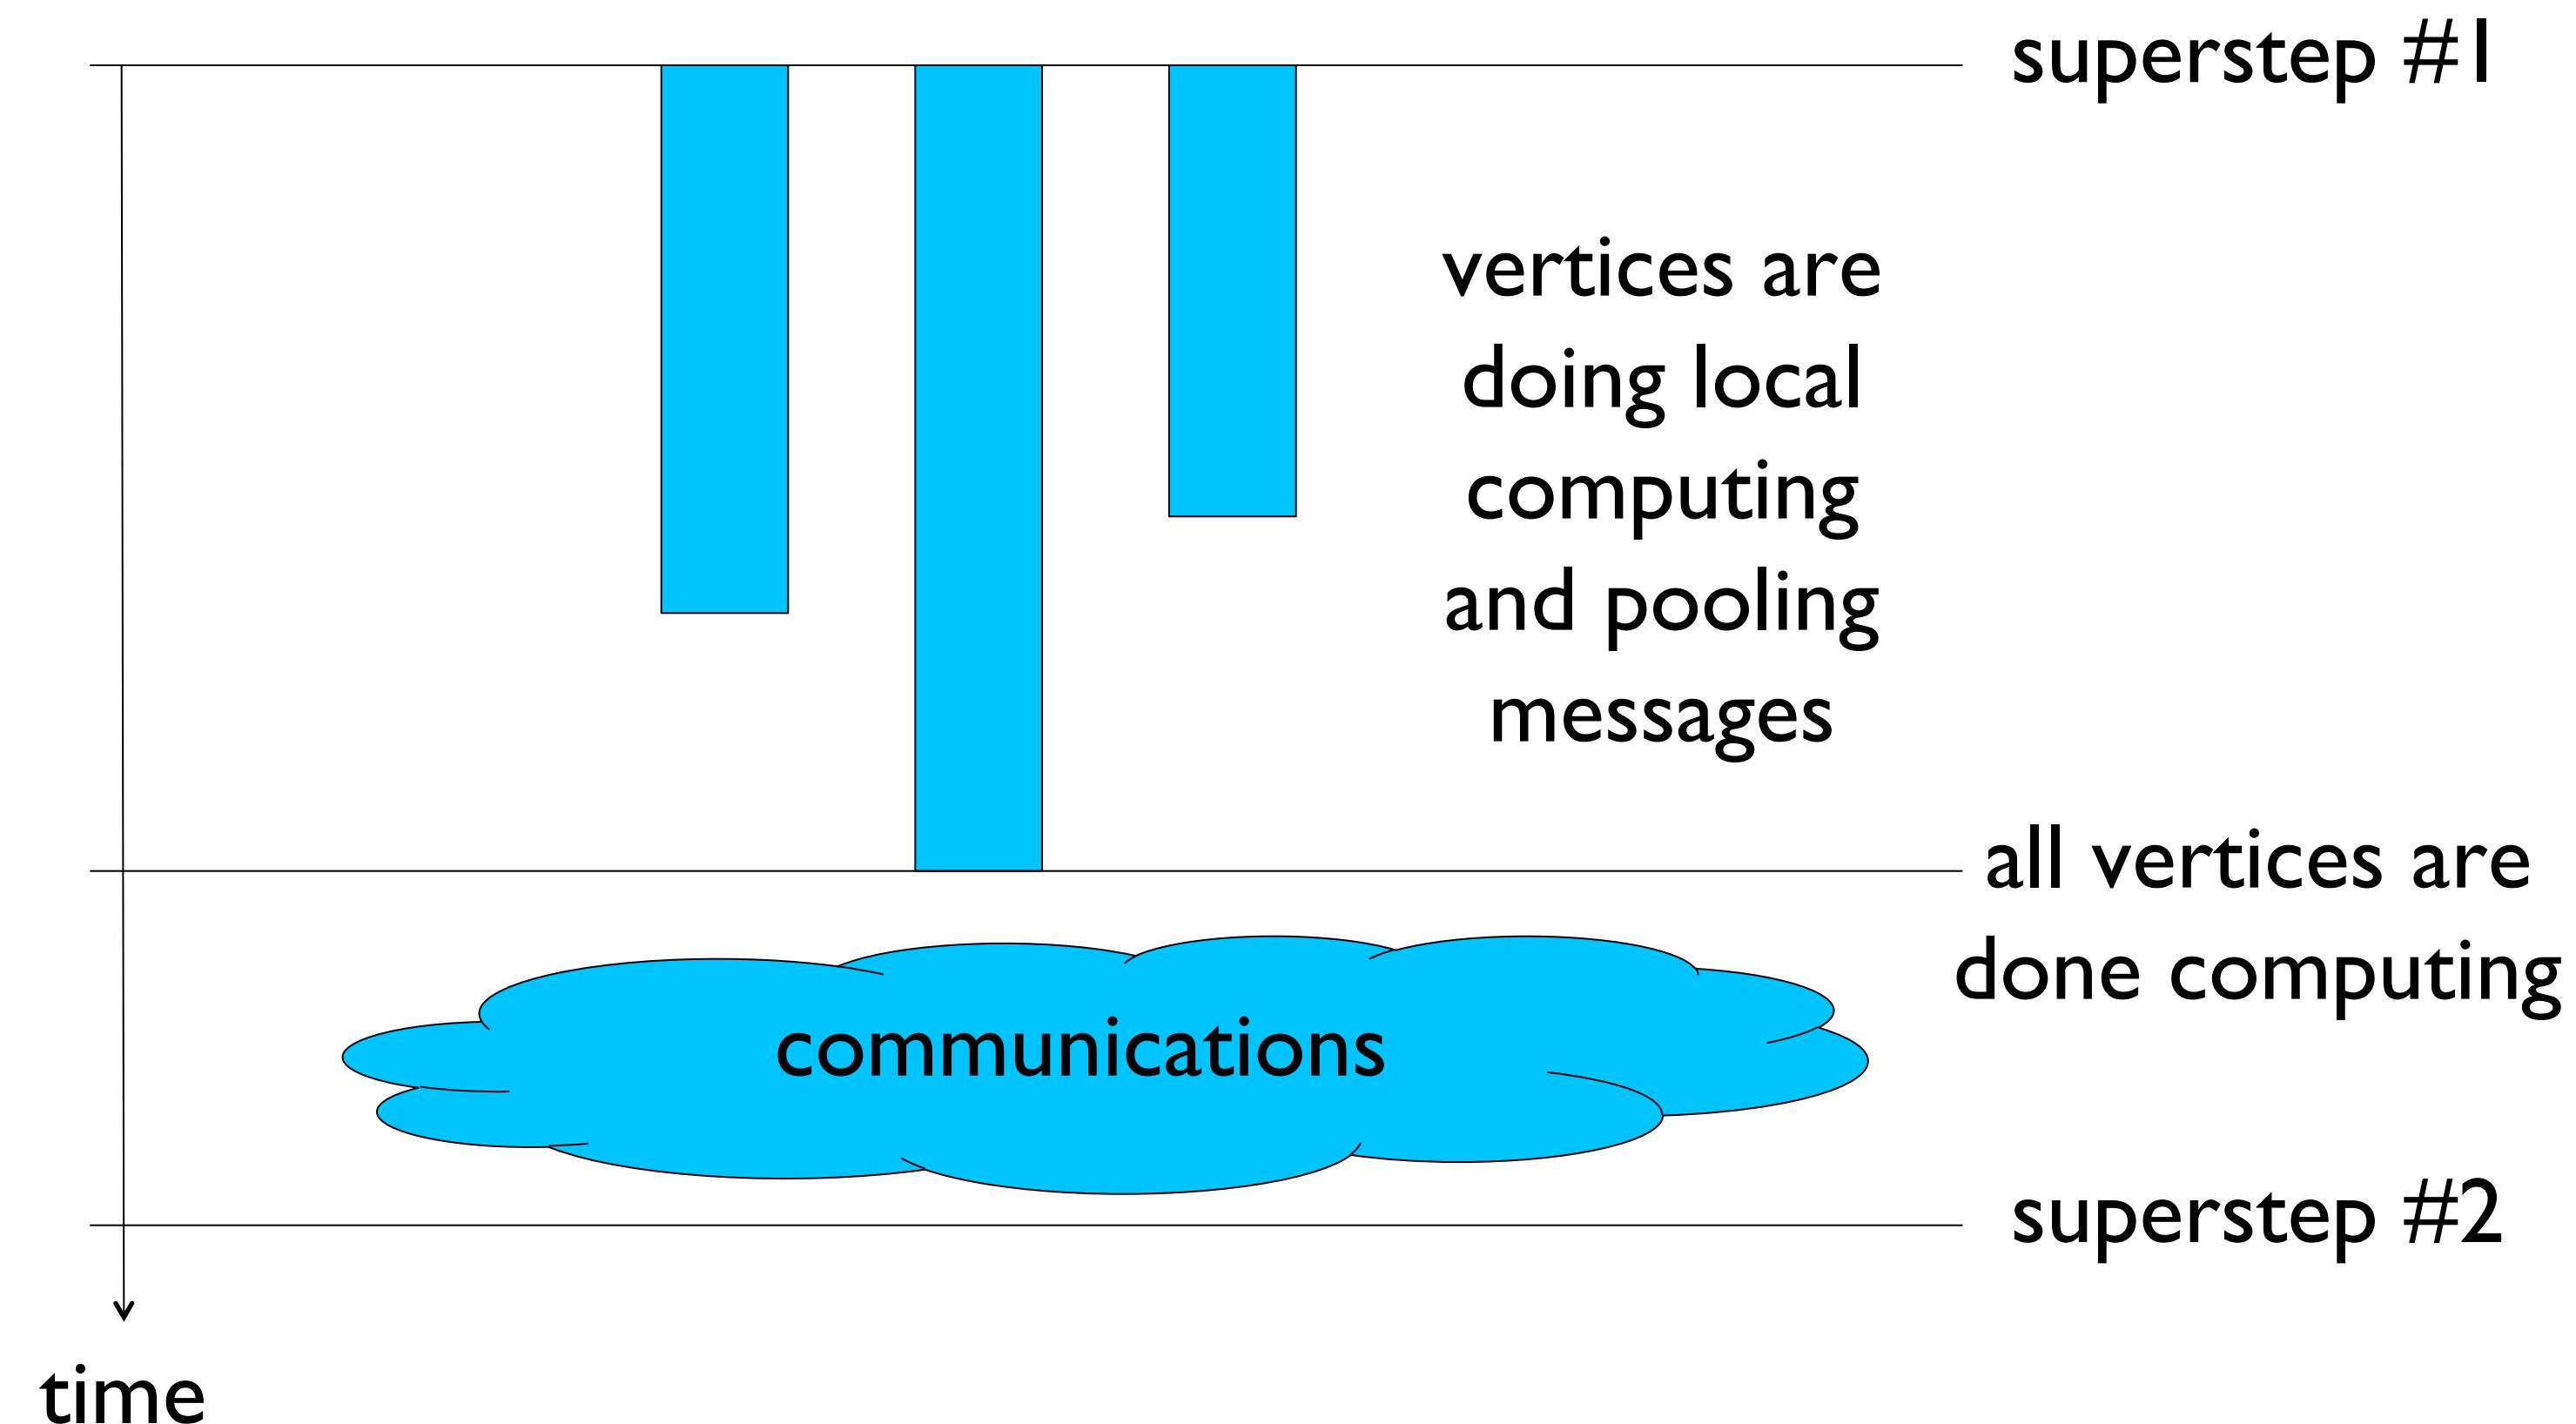
Vertex-centric BSP application



“Think like a vertex”

- I know my local state
- I know my neighbors
- I can send messages to vertices
- I can declare that I am done
- I can mutate graph topology

Local state, global messaging



O'REILLY®

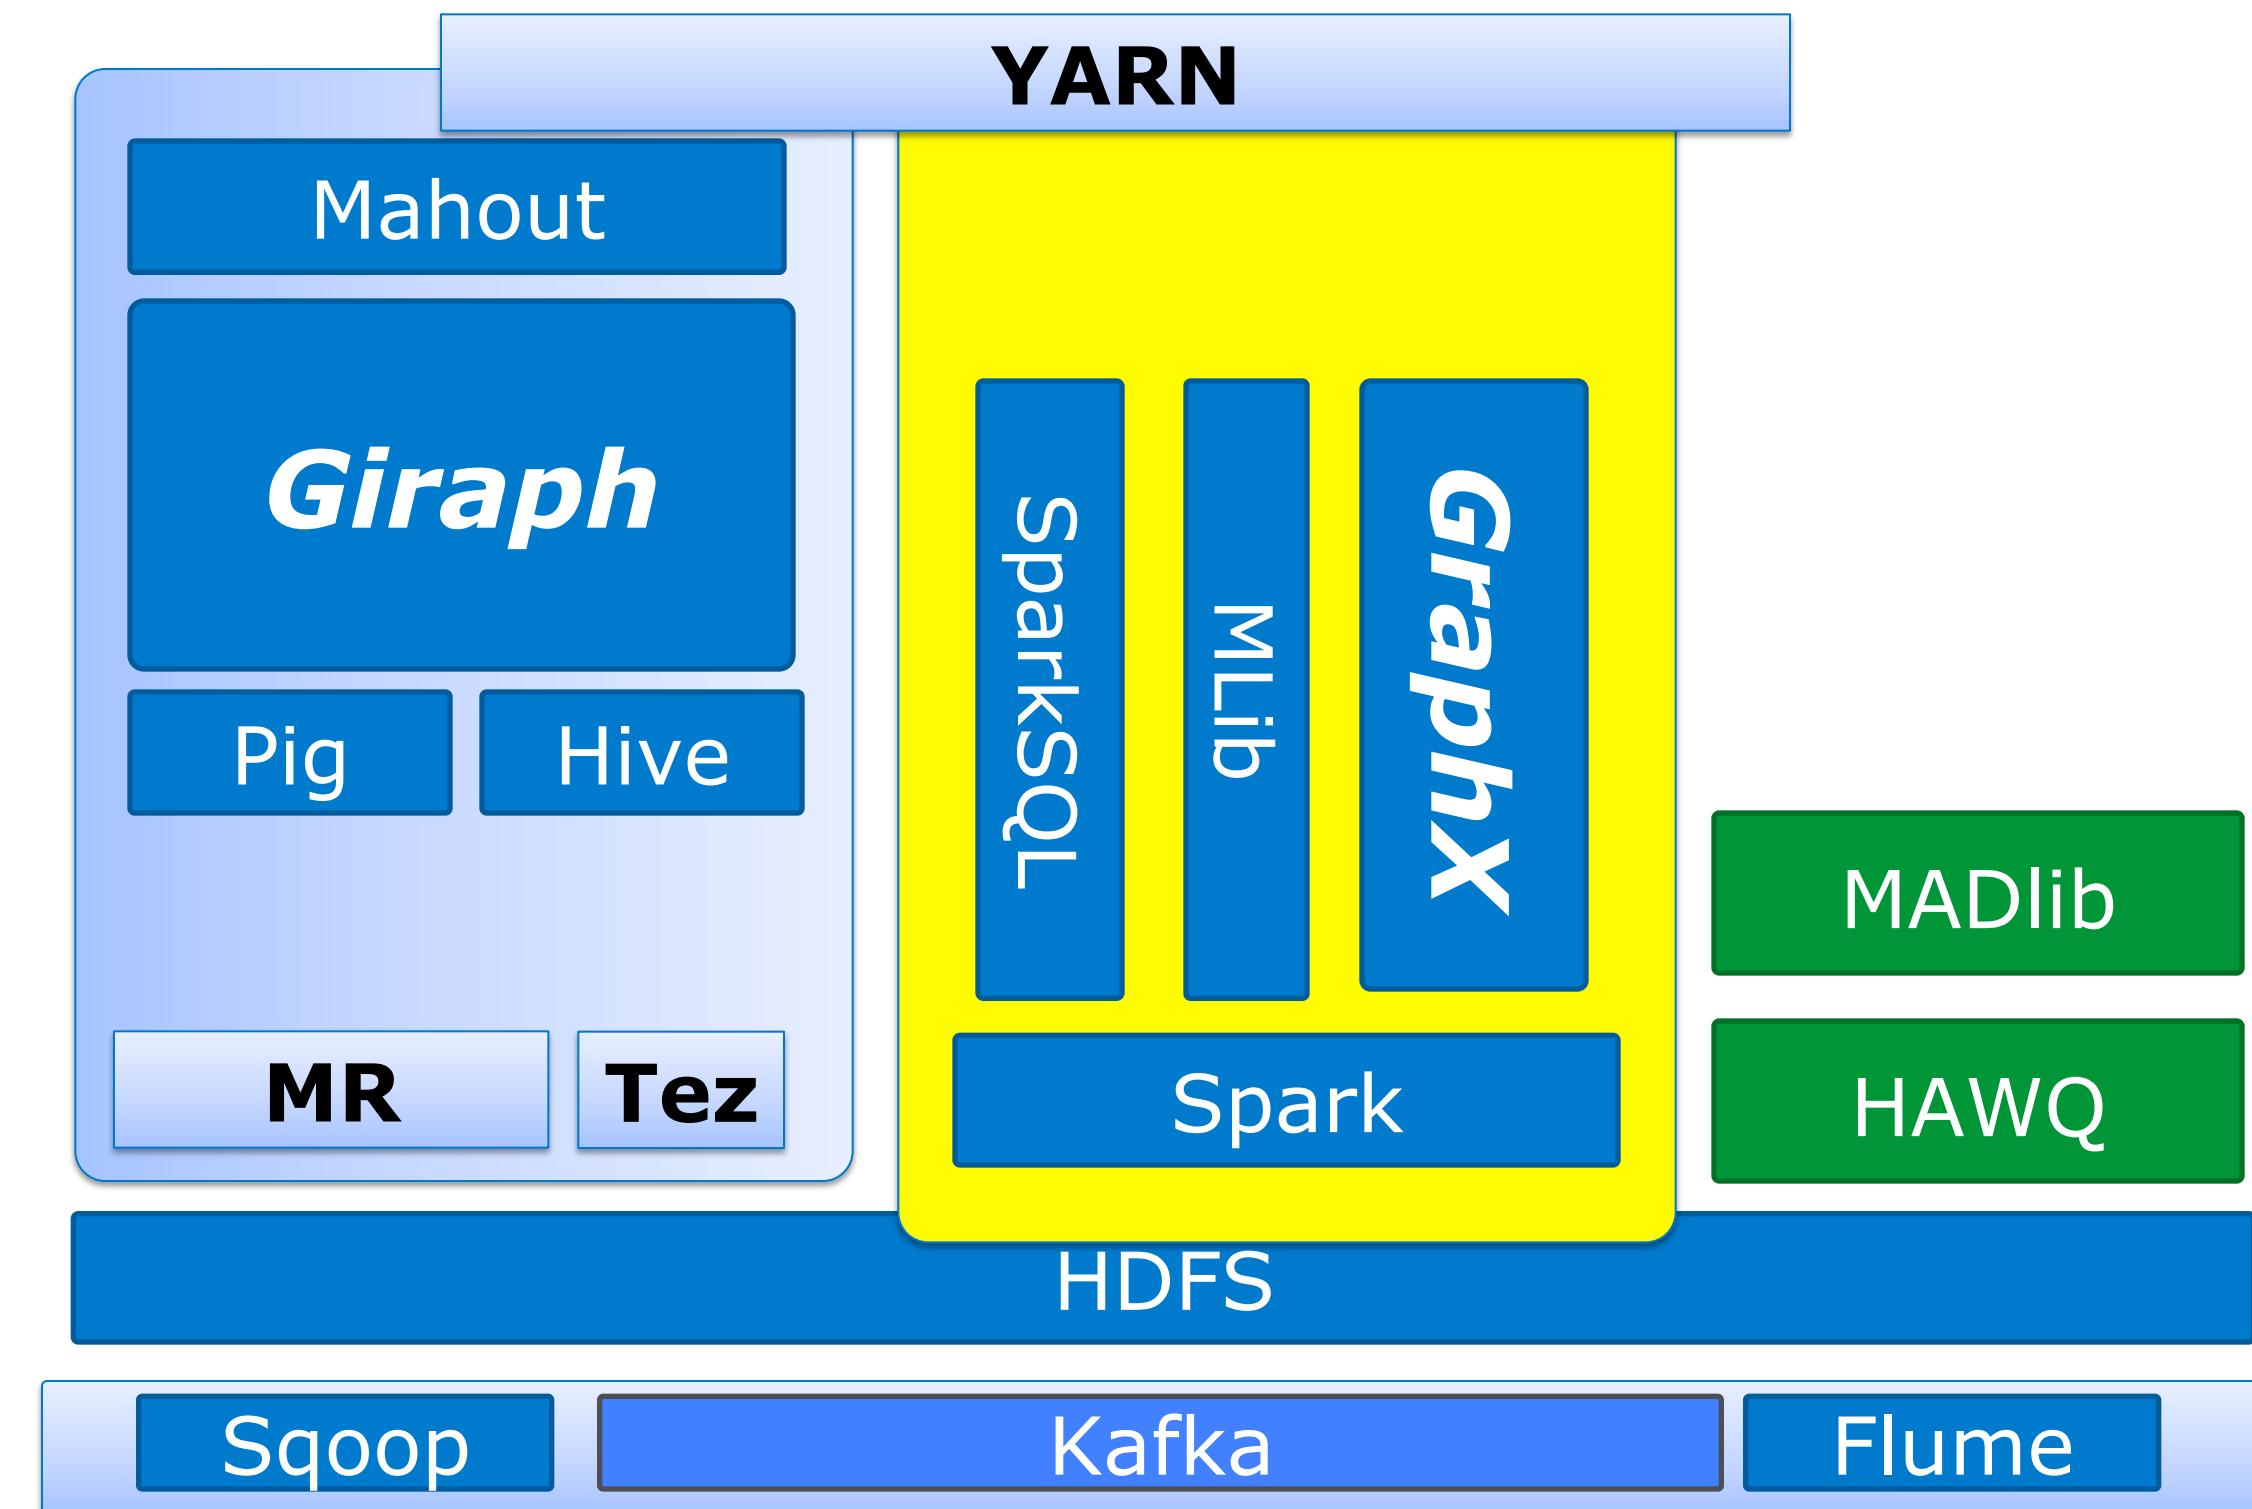
oscon

Lets put it all together

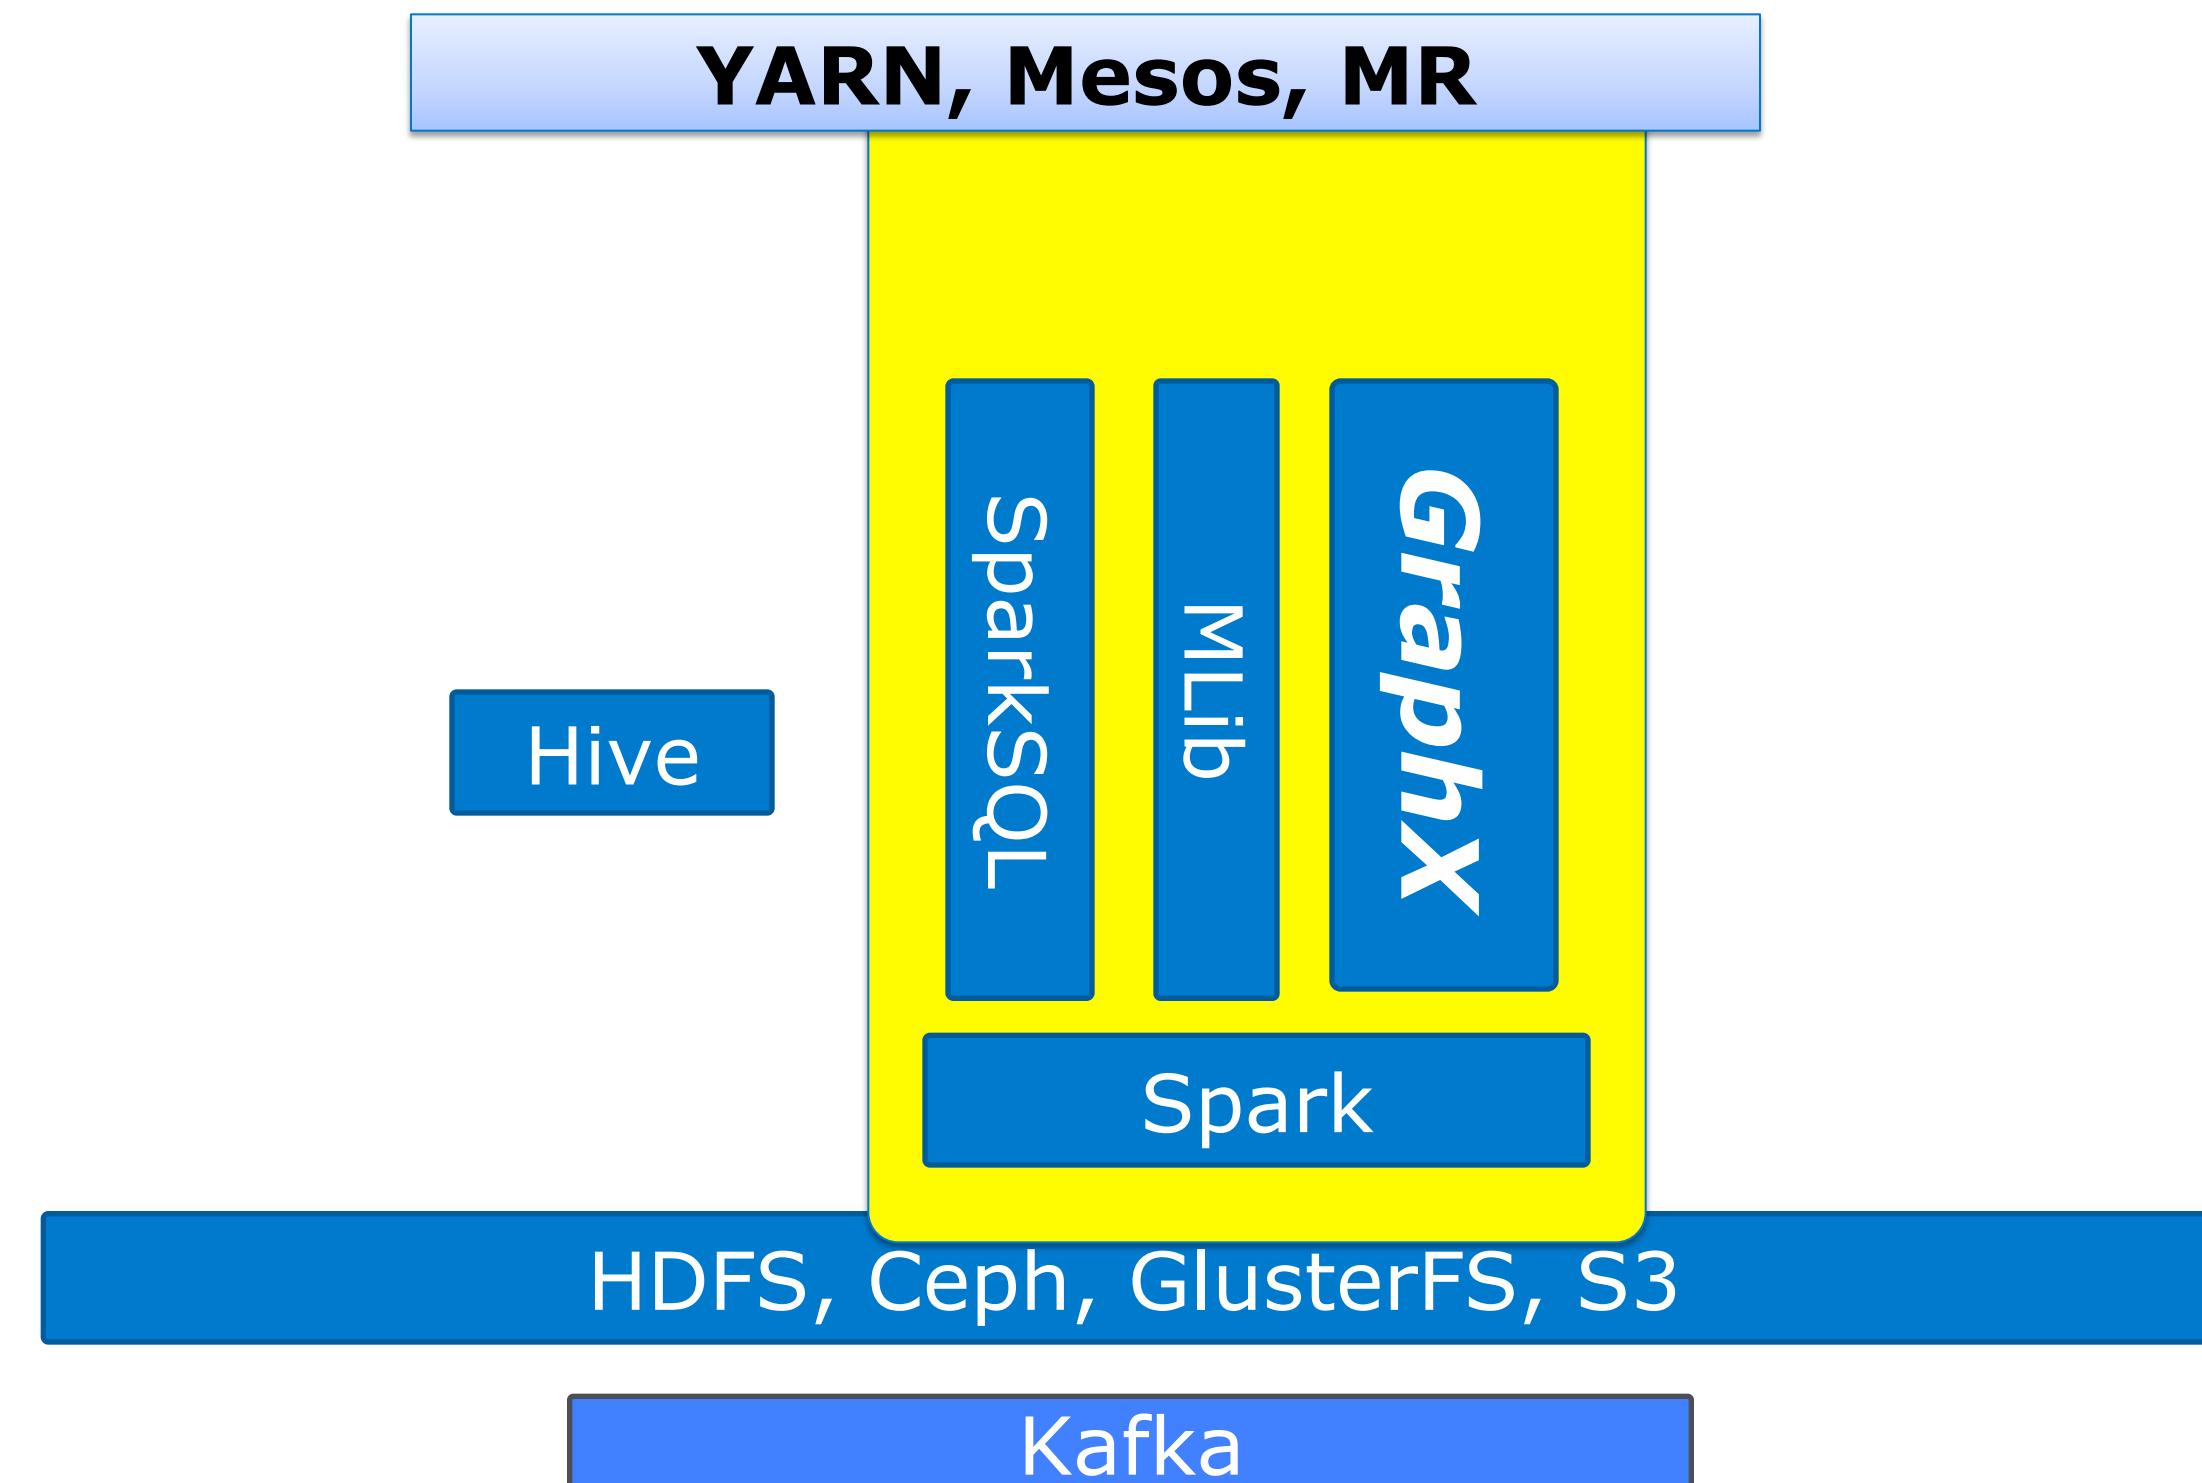
oscon.com

#oscon

Hadoop ecosystem view



Spark view



O'REILLY®

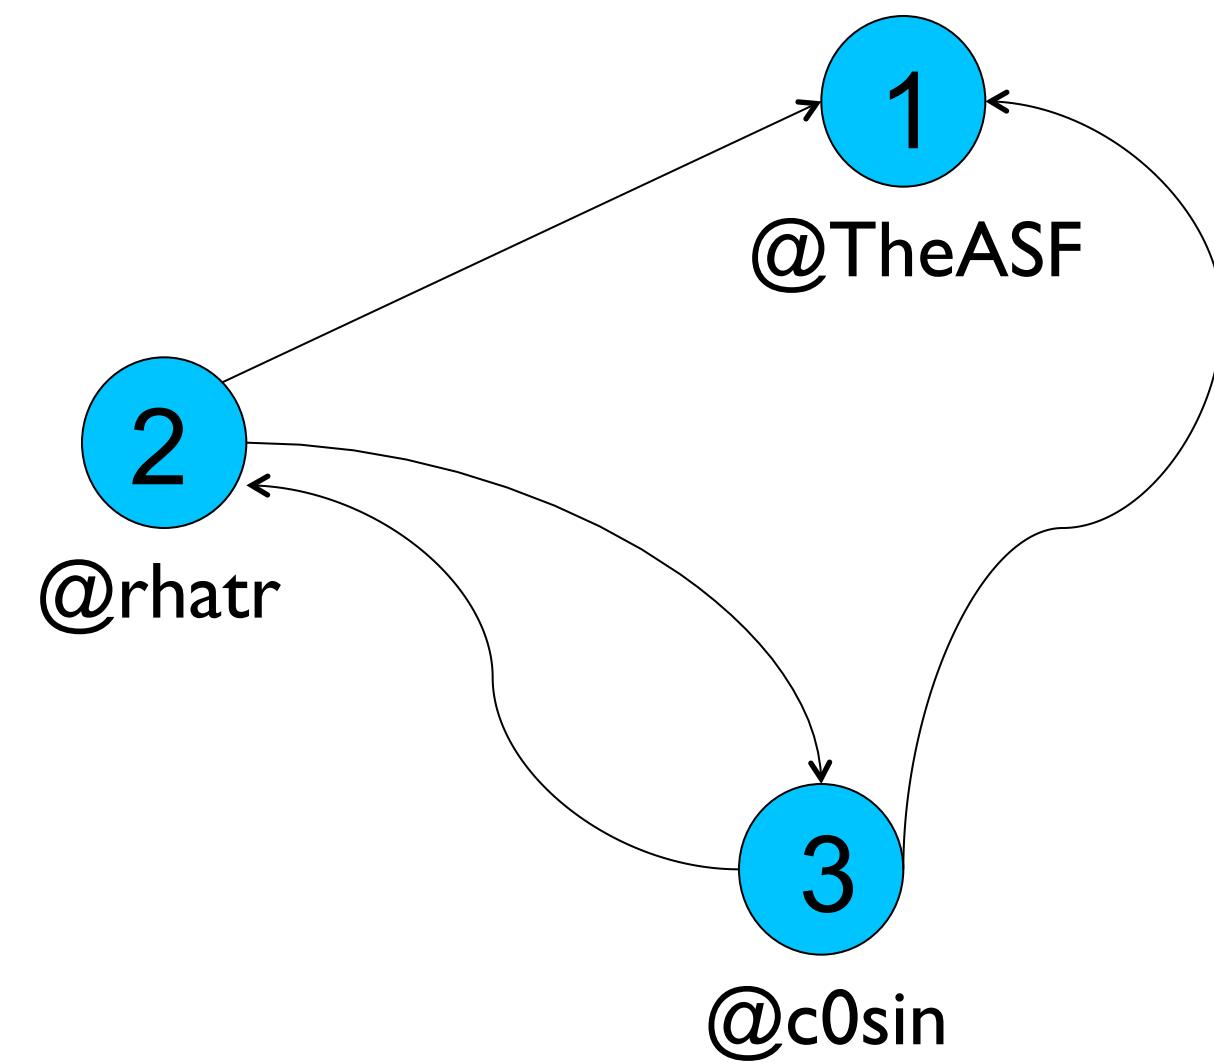
oscon

Enough boxology!
Let's look at some code

oscon.com

#oscon

Our toy for the rest of this talk



Adjacency lists stored on HDFS

```
$ hadoop fs -cat /tmp/graph/1.txt
```

```
1
```

```
2 1 3
```

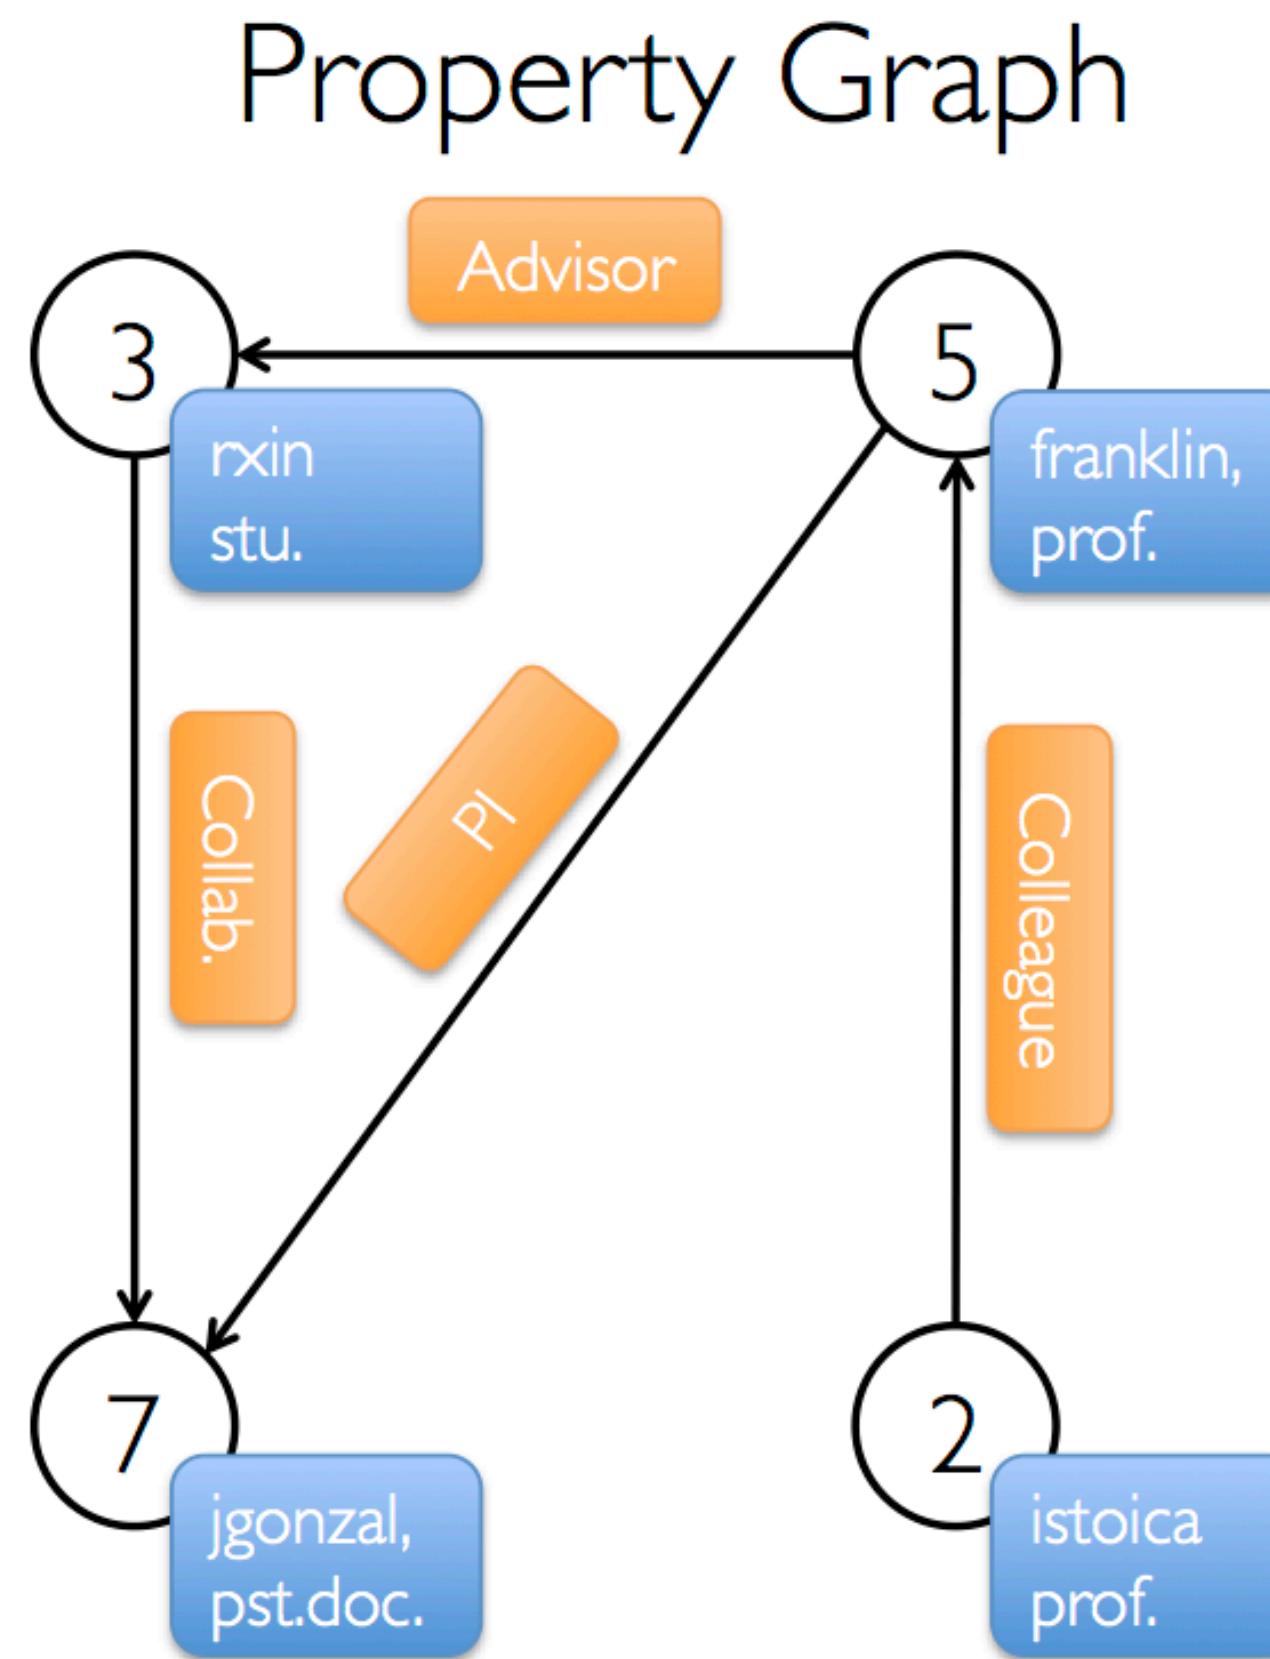
```
3 1 2
```

Graph modeling in GraphX

- The property graph is parameterized over the vertex (VD) and edge (ED) types

```
class Graph[VD, ED] {  
    val vertices: VertexRDD[VD]  
    val edges: EdgeRDD[ED]  
}
```

- Graph[(String, String), String]



Vertex Table

Id	Property (V)
3	(rxin, student)
7	(jgonzal, postdoc)
5	(franklin, professor)
2	(istoica, professor)

Edge Table

SrcId	DstId	Property (E)
3	7	Collaborator
5	3	Advisor
2	5	Colleague
5	7	PI

Hello world in GraphX

```
$ spark*/bin/spark-shell  
scala> val inputFile = sc.textFile("hdfs://tmp/graph/1.txt")  
scala> val edges = inputFile.flatMap(s => {  
    val l = s.split("\t");  
    l.drop(1).map(x => (l.head.toLong, x.toLong)) // [(2, 1), (2, 3)]  
})  
scala> val graph = Graph.fromEdgeTuples(edges, "") // Graph[String, Int]  
scala> val result = graph.collectNeighborhoods(EdgeDirection.Out).map(x =>  
    println("Hello world from the: " + x._1 + " : " + x._2.mkString(" ")))  
scala> result.collect() // don't try this @home
```

Hello world from the: 1 :

Hello world from the: 2 : 1 3

Hello world from the: 3 : 1 2

Graph modeling in Giraph

```
BasicComputation<I extends WritableComparable,  
                 V extends Writable,  
                 E extends Writable,  
                 M extends Writable>  
    // VertexID -- vertex ref  
    // VertexData -- a vertex datum  
    // EdgeData -- an edge label  
    // MessageData -- message payload
```

V is sort of like VD
E is sort of like ED

Hello world in Giraph

```
public class GiraphHelloWorld extends
    BasicComputation<IntWritable, IntWritable, NullWritable, NullWritable> {
    public void compute(Vertex<IntWritable, IntWritable, NullWritable> vertex,
                       Iterable<NullWritable> messages) {
        System.out.print("Hello world from the: " + vertex.getId() + " : ");
        for (Edge<IntWritable, NullWritable> e : vertex.getEdges()) {
            System.out.print(" " + e.getTargetVertexId());
        }
        System.out.println("");
        vertex.voteToHalt();
    }
}
```

How to run it

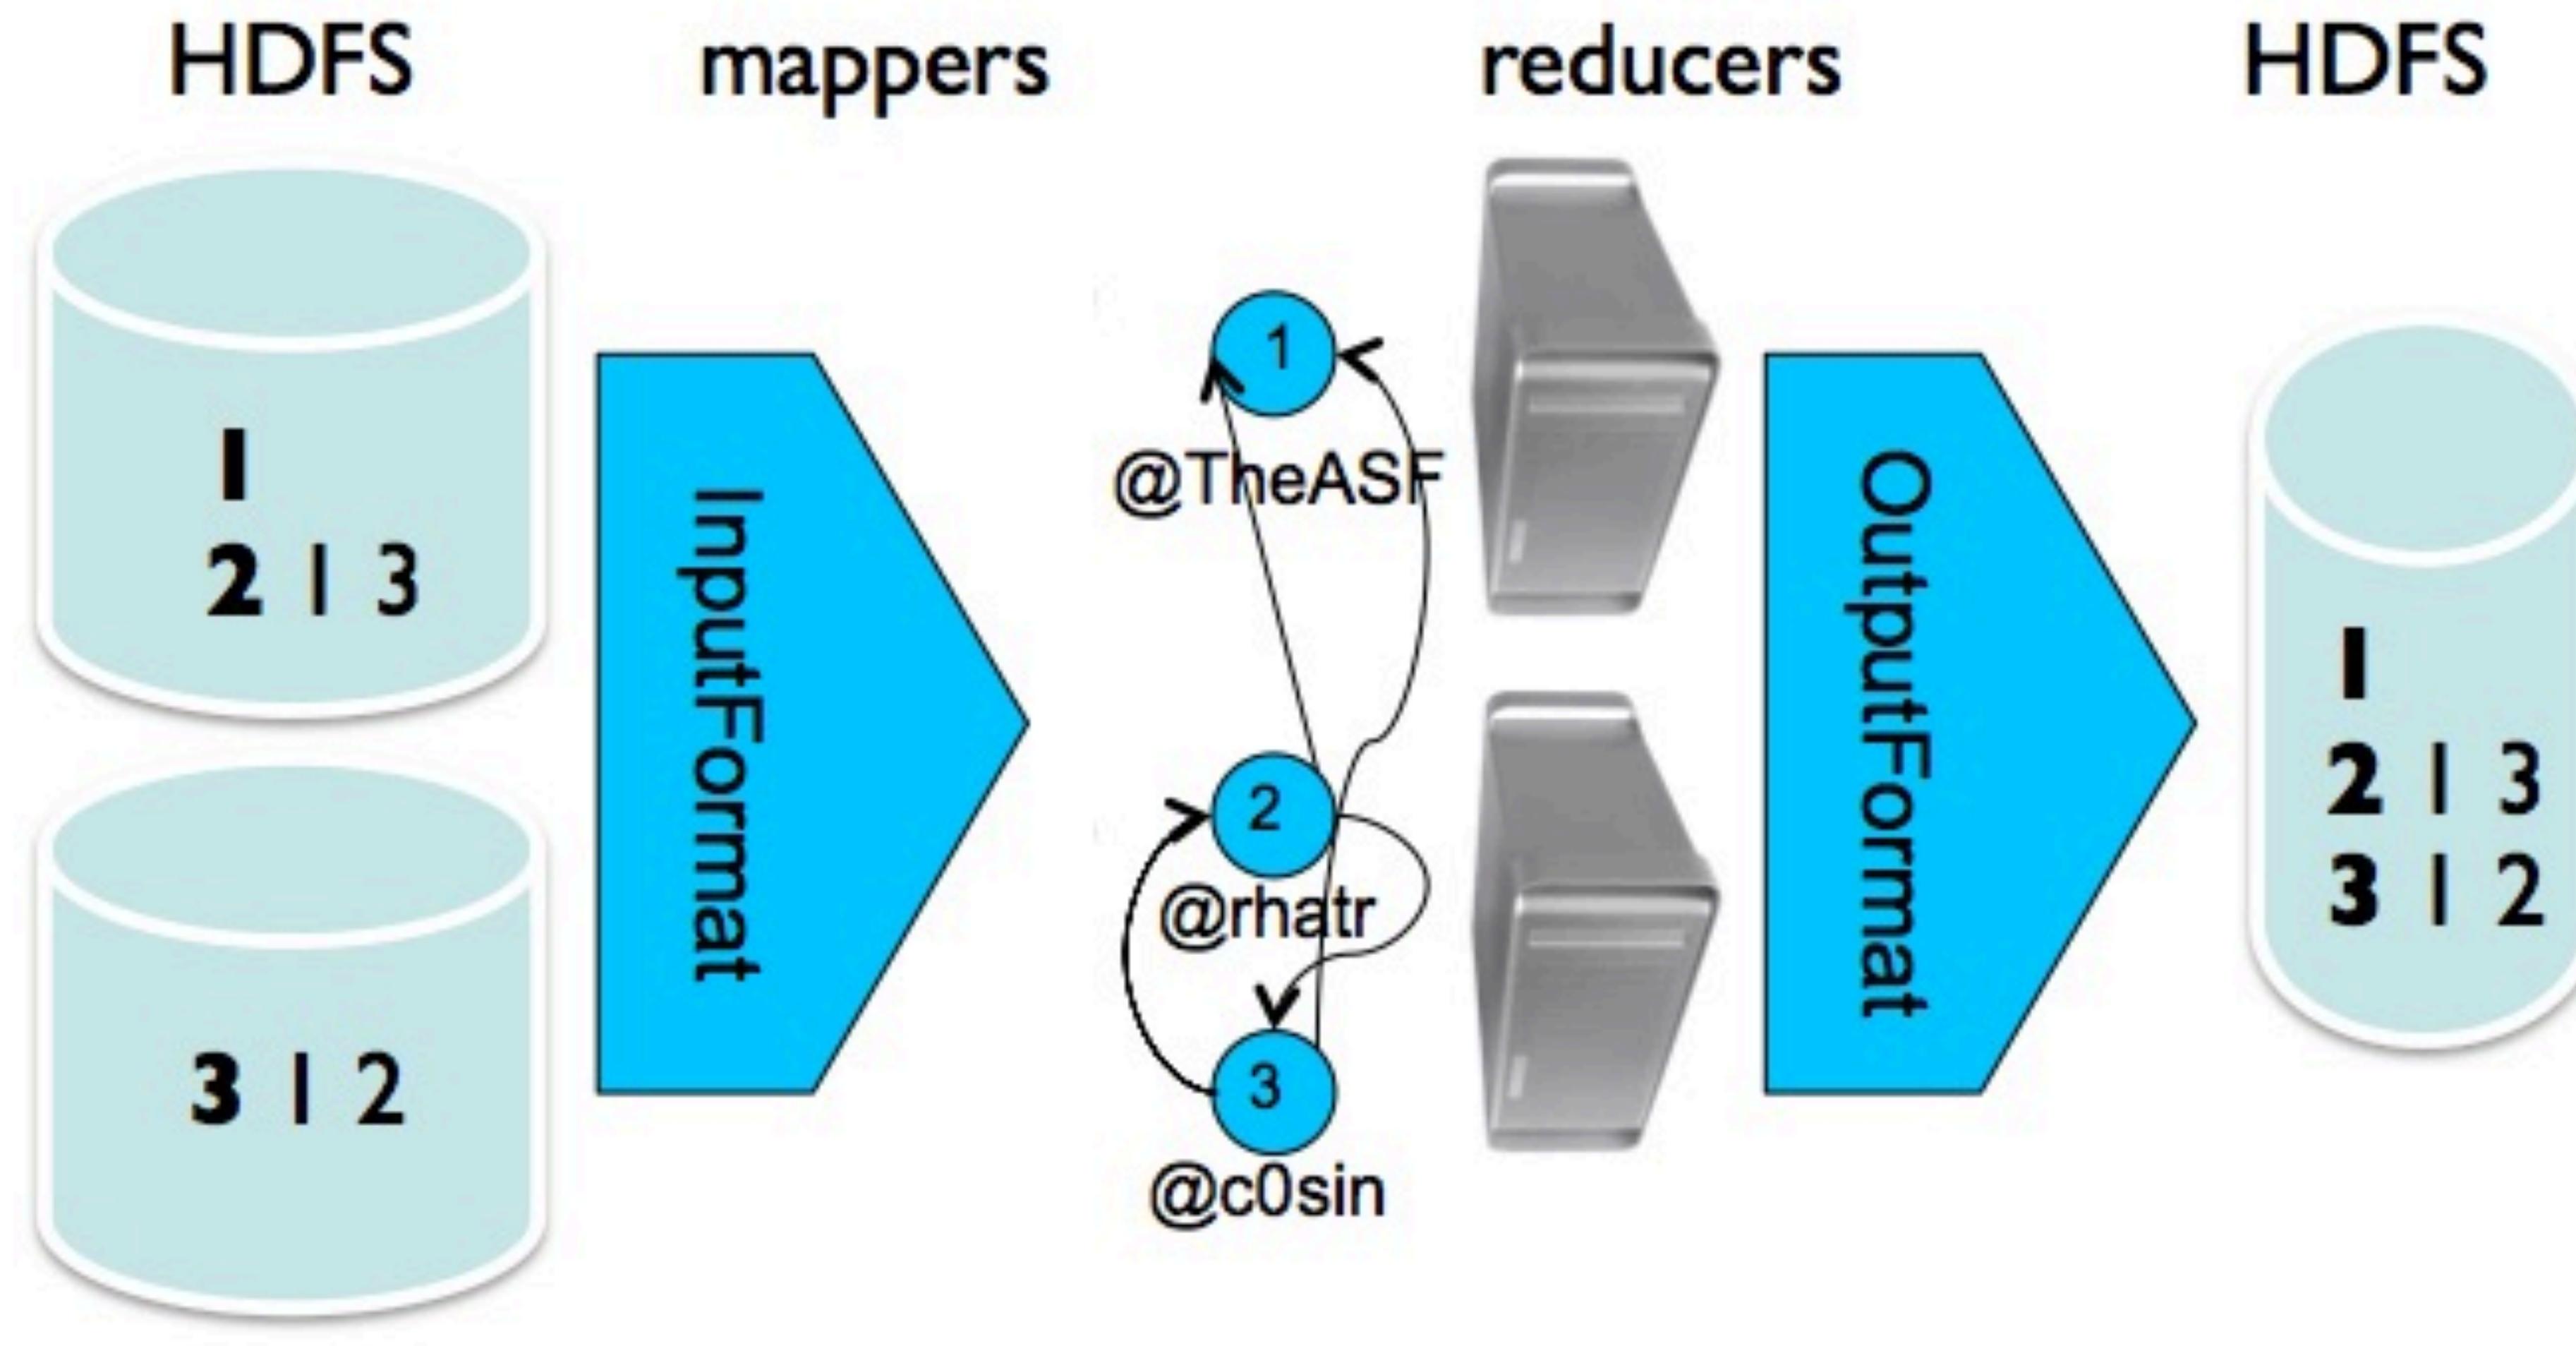
```
$ giraph target/*.jar giraph.GiraphHelloWorld \
-vip /tmp/graph/ \
-vif org.apache.giraph.io.formats.IntIntNullTextInputFormat \
-w 1 \
-ca giraph.SplitMasterWorker=false,giraph.logLevel=error
```

Hello world from the: 1 :

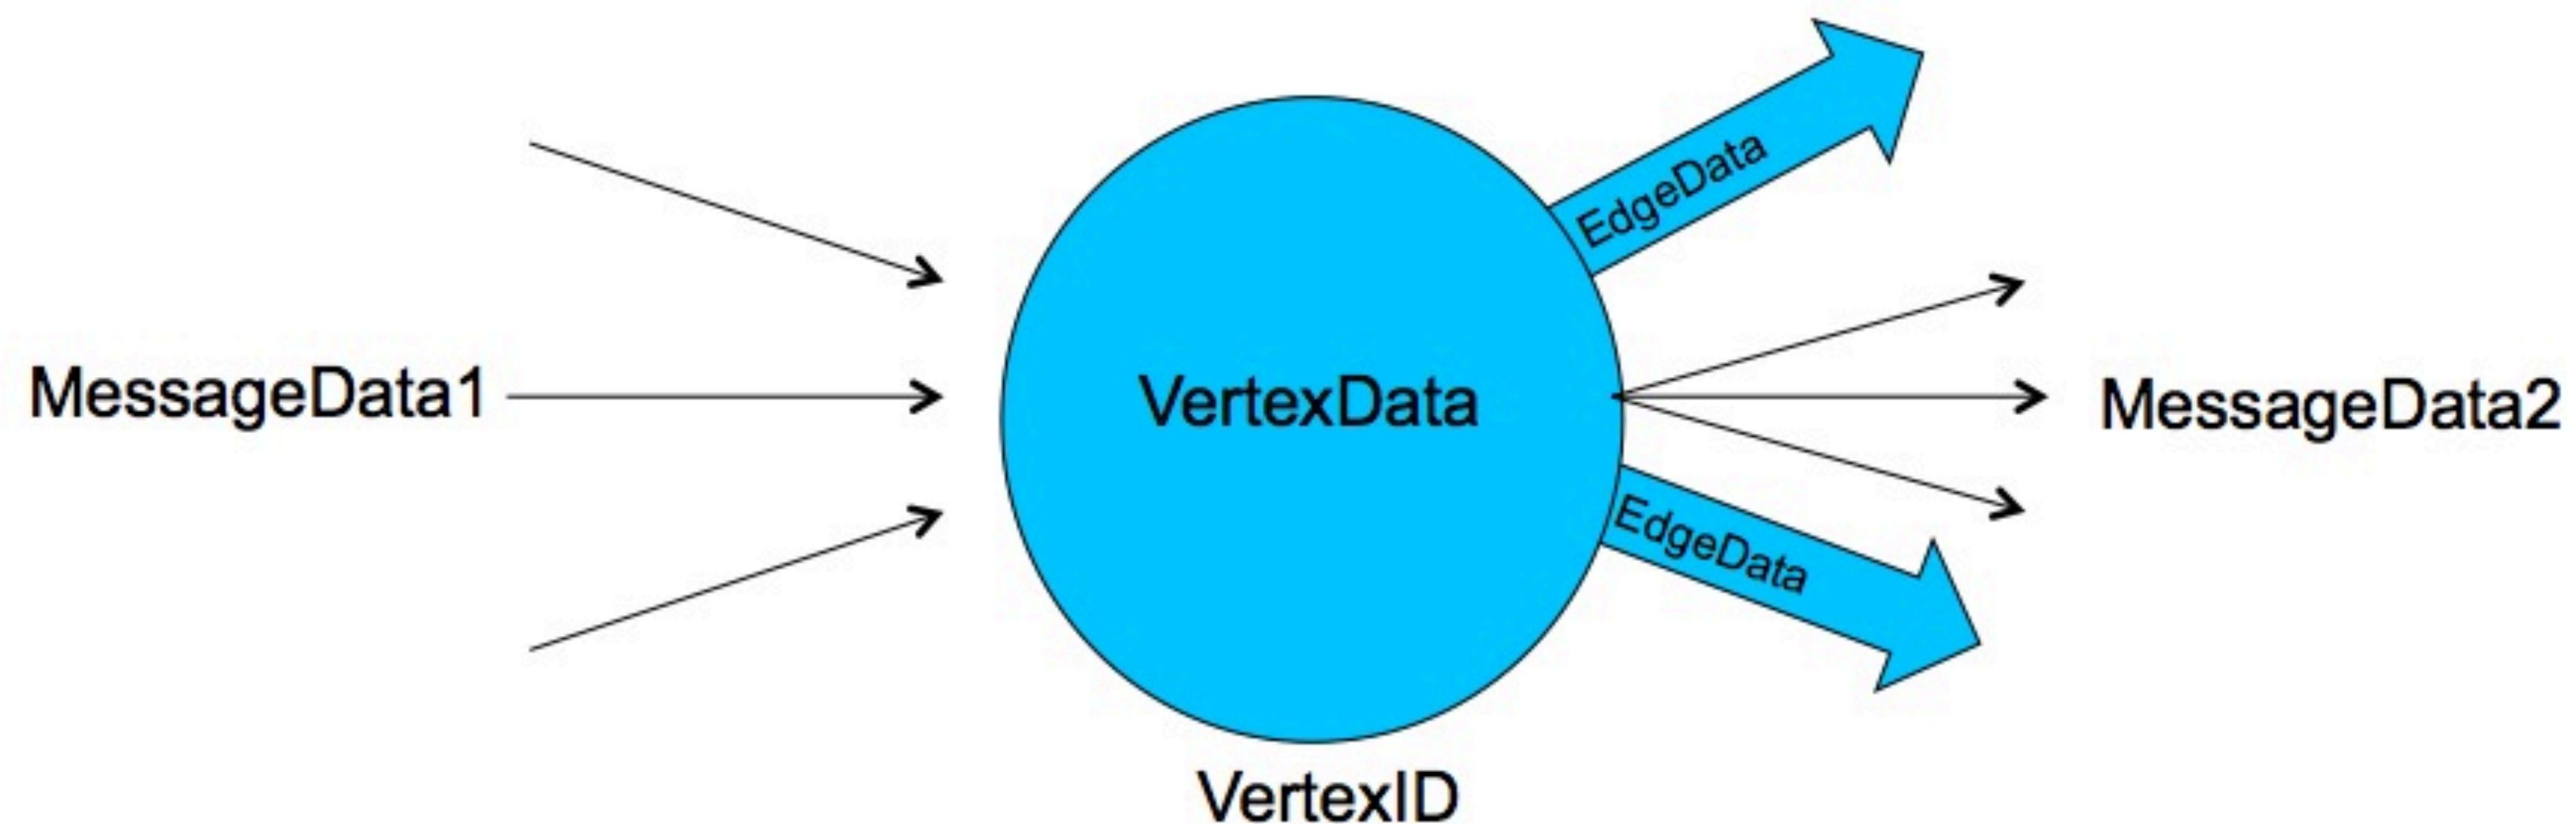
Hello world from the: 2 : 1 3

Hello world from the: 3 : 1 2

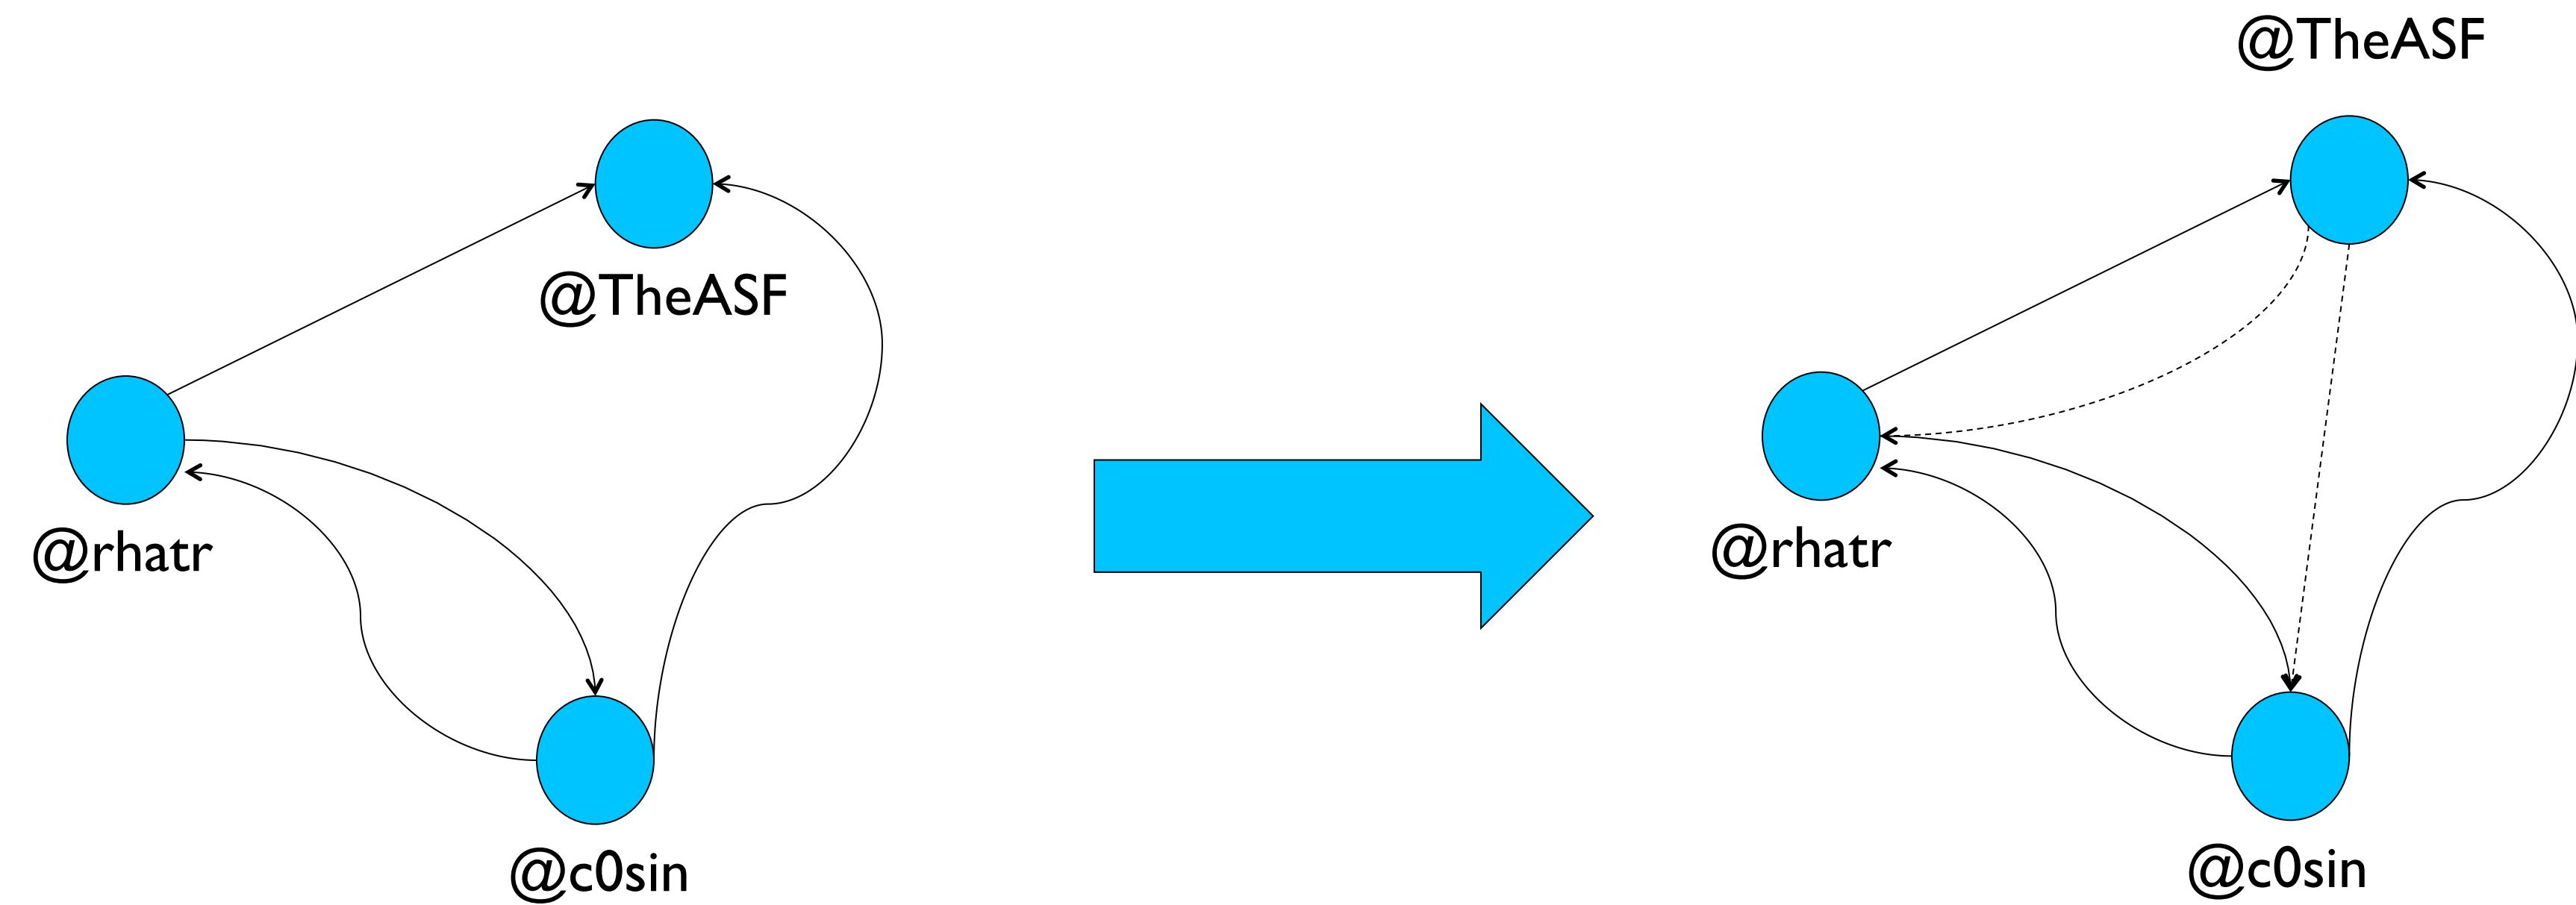
Anatomy of Giraph run



BSP assumes an exclusively vertex view



Turning Twitter into Facebook



Hello world in Giraph

```
public void compute(Vertex<Text, DoubleWritable, DoubleWritable> vertex, Iterable<Text> ms ) {
    if (getSuperstep() == 0) {
        sendMessageToAllEdges(vertex, vertex.getId());
    } else {
        for (Text m : ms) {
            if (vertex.getEdgeValue(m) == null) {
                vertex.addEdge(EdgeFactory.create(m, SYNTHETIC_EDGE));
            }
        }
    }
    vertex.voteToHalt();
}
```

O'REILLY®

oscon

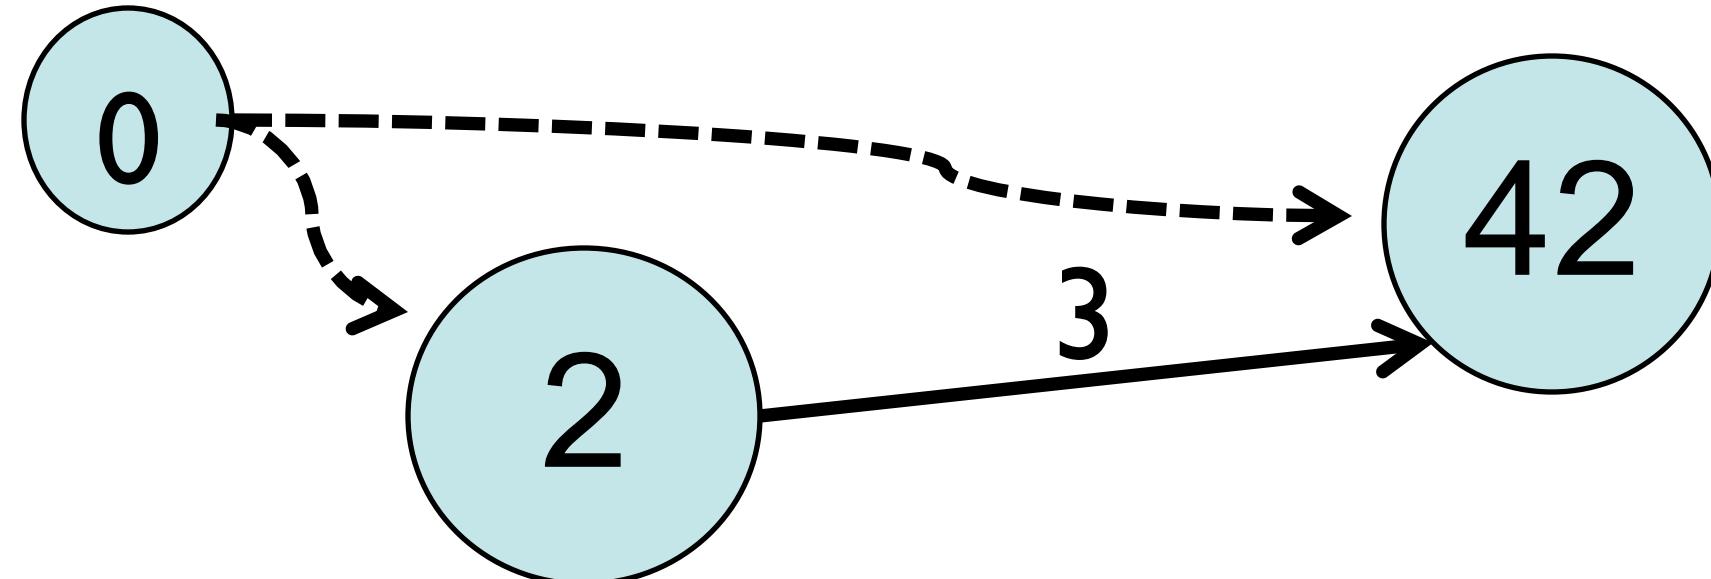
BSP in GraphX

oscon.com

#oscon

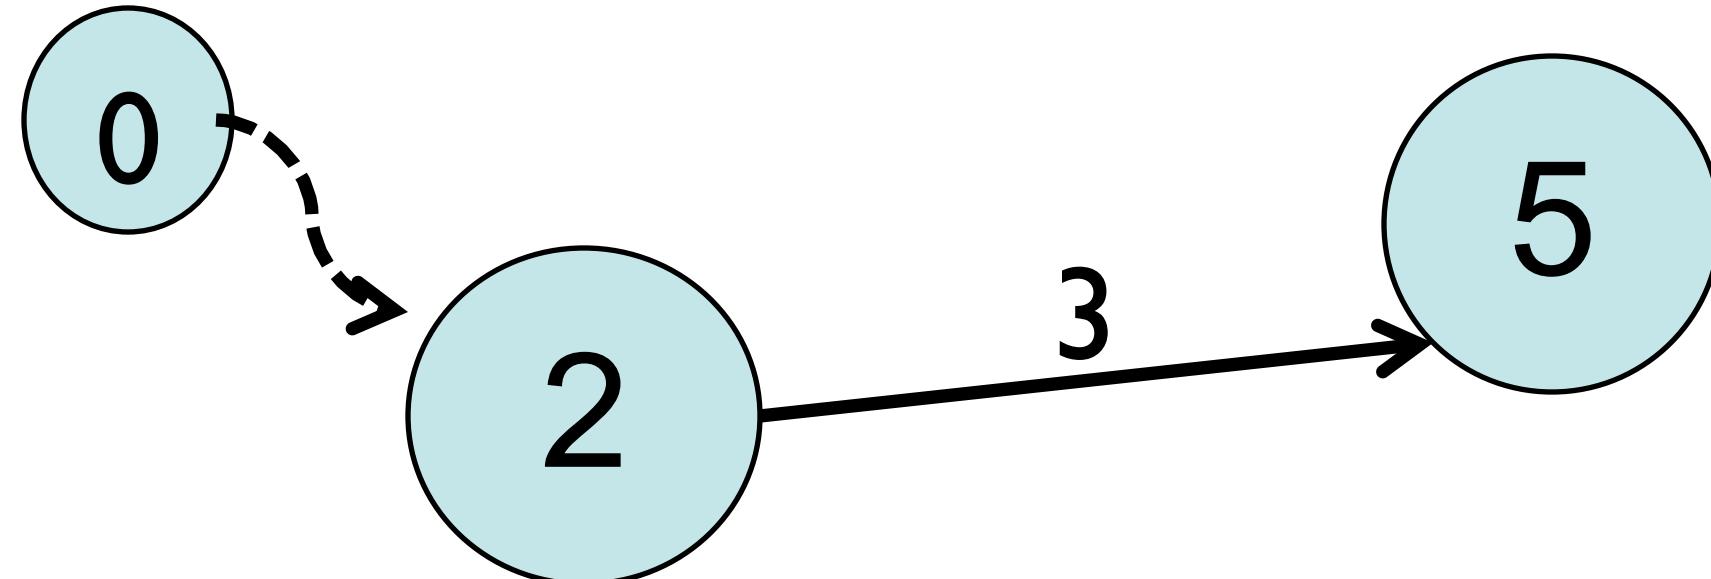
Single source shortest path

```
scala> val sssp = graph.pregel(Double.PositiveInfinity) // Initial message
((id, dist, newDist) => math.min(dist, newDist),           // Vertex Program
 triplet => {                                                 // Send Message
   if (triplet.srcAttr + triplet.attr < triplet.dstAttr) {
     Iterator((triplet.dstId, triplet.srcAttr + triplet.attr))
   } else {
     Iterator.empty
   }
},
(a,b) => math.min(a,b))                                     // Merge Messages
scala> println(sssp.vertices.collect.mkString("\n"))
```

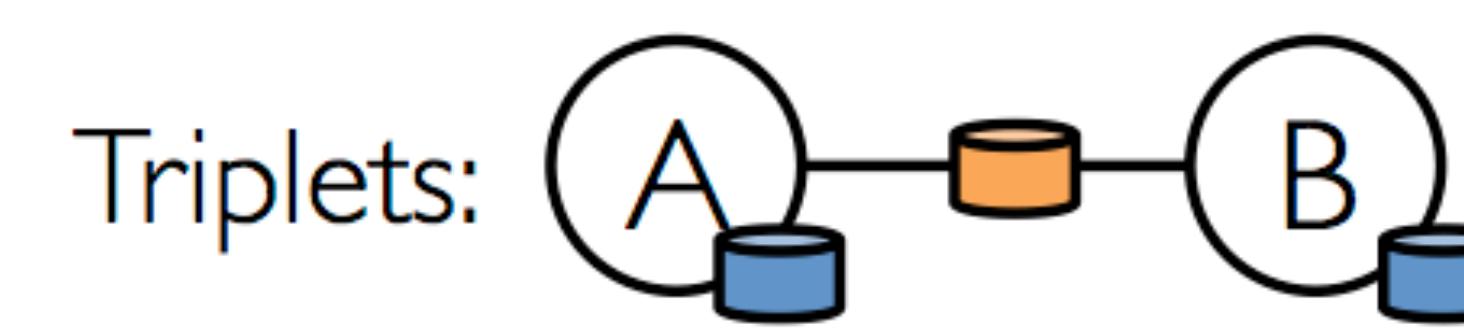
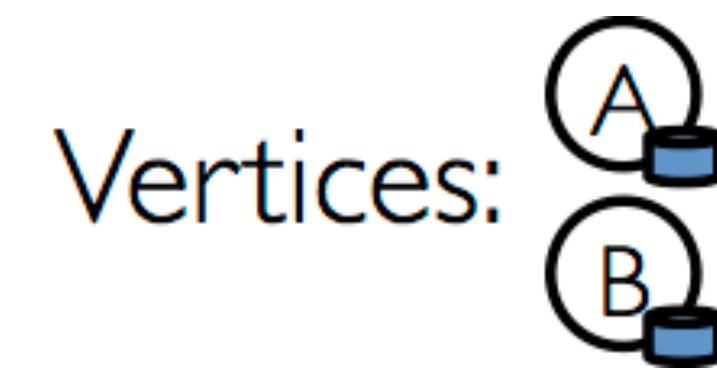


Single source shortest path

```
scala> val sssp = graph.pregel[Double.PositiveInfinity) // Initial message  
((id, dist, newDist) => math.min(dist, newDist), // Vertex Program  
triplet => { // Send Message  
    if (triplet.srcAttr + triplet.attr < triplet.dstAttr) {  
        Iterator((triplet.dstId, triplet.srcAttr + triplet.attr))  
    } else {  
        Iterator.empty  
    }  
},  
(a,b) => math.min(a,b)) // Merge Messages  
scala> println(sssp.vertices.collect.mkString("\n"))
```



Operational views of the graph



Masking instead of mutation

- def subgraph(
 epred: EdgeTriplet[VD,ED] => Boolean = (x => true),
 vpred: (VertexID, VD) => Boolean = ((v, d) => true))
 : Graph[VD, ED]
- def mask[VD2, ED2](other: Graph[VD2, ED2]): Graph[VD, ED]

Built-in algorithms

- `def pageRank(tol: Double, resetProb: Double = 0.15): Graph[Double, Double]`
- `def connectedComponents(): Graph[VertexID, ED]`
- `def triangleCount(): Graph[Int, ED]`
- `def stronglyConnectedComponents(numIter: Int): Graph[VertexID, ED]`

Final thoughts

Giraph

- An unconstrained BSP framework
- Specialized fully mutable, dynamically balanced in-memory graph representation
- Very procedural, vertex-centric programming model
- Genuine part of Hadoop ecosystem
- Definitely a 1.0

GraphX

- An RDD framework
- Graphs are “views” on RDDs and thus immutable
- Functional-like, “declarative” programming model
- Genuine part of Spark ecosystem
- Technically still an alpha

O'REILLY®

oscon

Q&A

Thanks!

oscon.com

#oscon